


<http://facilitiesci.org>

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

*Connecting Large Facilities,
Connecting CI, Connecting People*



The workshop was funded by the National Science Foundation under grant #1933353

September 16-17, 2019
Alexandria, VA

Disclaimer. This material is based upon work supported by the National Science Foundation under grant #1933353. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the workshop participants and do not necessarily reflect the views of the National Science Foundation.

Cover. Designed by Rafael Ferreira da Silva (University of Southern California). Cover images credits: The Ice Cube Neutrino Observatory¹, The Large Hadron Collider², The Nathaniel B. Palmer at McMurdo Station (Kelly Cheek, National Science Foundation)³, NEON Meteorological Flux Tower at terrestrial sites (Abraham Karam, NEON), Gray and Black Galaxy⁴, Aerial Photography of Wide Green Grass Field (Stephan Müller)⁵, Meeting photo (Facilities CI workshop)⁶, and Cable Network (Taylor Vick)⁷.

Preferred citation:

E. Deelman, I. Baldin, B. Bockelman, A. Bolton, P. Brady, T. Cheatham, L. Christopherson, R. Ferreira da Silva, T. Gulbransen, K. Keahay, M. Kogan, A. Mandal, A. Murillo, J. Nabrzyski, V. Pascucci, S. Petruzza, M. Rynge, S. Sons, D. Stanzione, C. Surajit, D. Swensen, A. Szalay, D. Thain, J. Towns, C. Vardeman, J. Wyngaard, “2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure: Connecting Large Facilities, Connecting CI, Connecting People”, 2020, DOI: <https://dx.doi.org/10.25549/0ZBF-8M77>.

¹ https://www.nsf.gov/news/mmg/media/images/ICL_SmallDomBlue_Horizontal.jpg

² <https://news.wisc.edu/content/uploads/2018/09/hadron-collider.jpg>

³ https://www.nsf.gov/news/mmg/media/images/NBP_H.jpg

⁴ <https://www.pexels.com/photo/sky-space-dark-galaxy-2150>

⁵ <https://www.pexels.com/photo/aerial-photography-of-wide-green-grass-field-753869>

⁶ <https://facilitiesci.github.io/2019/>

⁷ <https://unsplash.com/photos/M5tzZtFCOfs>

Report Authors:

Ewa Deelman (PI and Chair), University of Southern California

Ilya Baldin, RENCI, UNC - Chapel Hill

Brian Bockelman, Morgridge Institute

Adam Bolton, National Optical Astronomy Observatory

Patrick Brady, University of Wisconsin-Milwaukee

Tom Cheatham, University of Utah

Laura Christopherson, RENCI, UNC-Chapel Hill

Rafael Ferreira da Silva, USC Information Sciences Institute

Tom Gulbransen, Battelle, NEON

Kate Keahey, Argonne National Laboratory,

Marina Kogan, University of Utah

Anirban Mandal, RENCI, UNC-Chapel Hill

Angela Murillo, Indiana University-Purdue University Indianapolis

Jarek Nabrzyski, University of Notre Dame

Valerio Pascucci, University of Utah

Steve Petruzza, University of Utah

Mats Rynge, USC Information Sciences Institute

Susan Sons, Indiana University, CACR

Dan Stanzione, Texas Advanced Computing Center

Chaudhuri Surajit, Microsoft

Daryl Swenson, Oregon State University

Alexander Szalay, Johns Hopkins

Douglas Thain, University of Notre Dame

John Towns, NCSA

Charles Vardeman, University of Notre Dame

Jane Wyngaard, University of Notre Dame

Table of Content

EXECUTIVE SUMMARY	1
<i>Summary of the Workshop</i>	1
<i>Summary of the Discussions</i>	2
<i>Key Findings and Recommendations</i>	3
INTRODUCTION	5
<i>Overview and Goals</i>	5
<i>Workshop Organization and Attendees</i>	5
<i>Workshop Structure and Activities</i>	6
THEME I: IDENTIFYING CYBERINFRASTRUCTURE CHALLENGES THAT FACILITIES FACE	7
<i>Recommendations</i>	9
THEME II: EXPLORING THE OPPORTUNITIES AND OBSTACLES TO COLLABORATION BETWEEN LFs	10
<i>Recommendations</i>	12
THEME III: EXAMINING NON-TECHNICAL CHALLENGES THAT INFLUENCE CI DEVELOPMENT	13
<i>Recommendations</i>	14
THEME IV: DEVELOPING IDEAS FOR ENHANCING THE CI WORKFORCE AND BUILDING A COMMUNITY OF CI PROFESSIONALS	14
<i>Recommendations</i>	16
CYBERINFRASTRUCTURE CALLING CARD SUMMARY	17
PARTICIPANT AND PRACTITIONER SURVEYS SUMMARY	19
<i>Workshop Participant Survey</i>	19
<i>Cyberinfrastructure Practitioner Survey</i>	22
REFERENCES	25
APPENDIX A: WORKSHOP CONTRIBUTORS	26
APPENDIX B: AGENDA	30
APPENDIX C: CYBERINFRASTRUCTURE CALLING CARDS.....	33



Executive Summary

Summary of the Workshop

In 2015, the National Science Foundation (NSF) began to support a series of biennial workshops that bring together large facilities (LFs) and cyberinfrastructure (CI) projects to share common experiences and challenges, discuss potential collaborations, and identify opportunities for leveraging CI within the community. The 2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure aimed to continue and advance this discussion, enable the exchange of CI solutions and challenges, and foster CI community building.

With a theme of “Connecting Large Facilities, Connecting CI, Connecting People,” the 2019 workshop emphasized the need to facilitate collaborations among LF and CI projects and recognized the importance of the CI workforce to LF science missions. Specific goals of the workshop included: 1) identifying CI challenges that facilities face when supporting their science missions; 2) exploring the opportunities and obstacles to collaboration between LFs; 3) examining non-technical challenges that influence CI development; and 4) developing ideas for enhancing the CI workforce and building a community of CI professionals.

A steering committee of cyberinfrastructure experts from NSF LFs and selected CI projects defined the workshop agenda with input gathered through a pre-workshop survey. Close to 100 representatives from NSF LFs, major NSF-funded CI projects, selected cloud providers, and NSF staff participated in presentations, panels, and breakout sessions centered around the workshop goals.

Prior to the workshop, a survey was sent to the participants to gauge their willingness to share CI resource and expertise and to share workforce development efforts. The 32 survey responders expressed ambivalence with regard to sharing resources but enthusiasm at the prospect of sharing expertise. The most frequently cited obstacles to developing a shared services effort were limited funding and personnel as well as divergent missions and goals of individual LFs. When asked about their willingness to contribute to a shared workforce development effort, most respondents indicated that they did not have resources to put toward this type of activity. Obstacles to sharing in the area of workforce mirrored those to sharing resources.

In September 2019, to obtain a quick pulse of the community, workshop participants were asked to distribute a survey to their teams to gain information about CI knowledge sharing practices, incentives to following a CI career path, and feelings of belonging to the broader professional CI community. Seventy-eight respondents completed this survey. The survey revealed that the LF community relies on online sources for information seeking and knowledge sharing. Key findings in the area of job satisfaction were that the CI practitioners are motivated primarily by the “sense of purpose/helping science/the world” followed closely by the ability to conduct research and solve complex problems. Although many respondents felt a belonging to their particular community, LF, or project, the survey did not show broader community connections.

A community building effort was also launched by requesting workshop participants to provide “CI calling cards” with professional information, pictures, and responses to questions about recent CI frustrations and successes. Frustrations revolved around data management challenges, budgetary constraints, communication (e.g., changing mindsets), workforce (e.g., hiring and retention), technical aspects of LF work (e.g., planning and design, the tension between older existing and emerging technologies, and juggling heterogeneous CI). Successes included developing new systems or services, improving existing ones, continuing to do good work throughout the year, and bringing people together via effective communication. These calling cards greatly facilitated group and one-on-one discussions.



Workshop materials, including the presentations and calling cards, are available on the workshop website <http://facilitiesci.org>.

Summary of the Discussions

LF cyberinfrastructure challenges

LFs face a spectrum of CI challenges when supporting their science missions. They collect and disseminate very large, diverse, and heterogeneous data sets with different temporal and spatial scales. At the same time, CI and commodity technologies are rapidly evolving. LF mission focus and the need to mitigate long-term risk often mean that software is developed in-house rather than adopted from existing solutions or created through collaborations. Timelines and budgets also may drive LFs to pursue expedient solutions rather than examine potential new designs and technologies. The same pressures can limit LF willingness to seek and maintain collaborations that could increase CI reuse and adoption.

In addition to the technical challenges, LFs face ever-growing demands in the area of workforce training, development, and retention. They are constantly competing with industry to attract and retain CI talent and must often hire domain scientists, programmers, or administrators without the needed experience to navigate the complex CI landscape. Even as the CI workforce matures, quickly changing technology makes it difficult for individual CI practitioners and LFs to keep up with new capabilities.

Opportunities and obstacles to collaboration between LFs

LFs have limited resources to collaborate with each other and large CI projects as they must focus on delivering instruments, data, and software to their communities. Although collaborations are potentially desirable and beneficial, it takes time to discover appropriate partnerships, understand common interests, and maintain a collaboration until it starts bearing fruit. The nature and extent of CI collaboration may vary across the spectrums of CI developers' expertise in computer science and scientific domains. There are also often mismatches between LF construction timelines and CI projects. Finally, LFs frequently perceive risks with engaging in external collaborations during the construction phase, especially if the projects and programs do not have common NSF program officers.

As CI complexity increases, there is a growing need for sharing communication, coordination, information, and expertise across the facilities. The workshop identified simple sharing approaches that could benefit LFs in the short term. For example, sharing knowledge of existing LF architectures can inform new LF design, and CI project training opportunities could be shared with LF staff. Guidance on open source licenses, community data standards, workflows, and software development best practices could also help LFs develop interoperable CI components and more easily discover potentially useful CI services. As collaborations and trust amongst LFs and the broader CI community grows, it may be possible to foster more tightly integrated activities such as co-development benefitting shared interests and the evaluation of new technologies.

Non-technical challenges influencing CI development

The non-technical challenges discussions focused on the tension between the resources available to support CI and the need to prioritize activities under these constraints. Often, CI improvements appear to receive less attention and associated investment compared to primary science instruments. Additionally, funding agencies frequently prioritize new software development over sustaining and maintaining existing CI, which increases the number of CI migration cycles. These conditions leave LFs having to choose between



innovating, upgrading capabilities, and maintaining current solutions while still delivering products and enabling science.

Approaches to enhancing CI workforce and building a community of CI professionals.

Discussions centered around CI workforce were extensive and permeated several workshop sessions. CI practitioners are a mix of domain- and computer-science educated personnel who often need extensive generic and facility-specific CI training. Computer science graduates can be hard to retain in a competitive market with higher salaries in non-science industries. Additionally, their skillsets often do not exactly match LF needs. With no curricula or certifications in the area of CI, finding or even describing the skills needed is challenging. Hiring is also complicated by the fact that enterprise skill sets do not directly translate to the CI environment and some LFs operate in remote locations. To improve hiring and retention, more emphasis is needed on LF career advantages, such as the opportunity to take part in an intrepid science endeavor and better work/life balance than many industry positions.

Diverse backgrounds, siloed CI development, and deployment processes within LFs make it challenging to build a broader CI community. There is no shared description nor clear definition of the skills, interests, and career paths for CI practitioners. Although communication channels exist within each LF, it is hard to catalyze communication and foster community across LFs and CI projects.

Key Findings and Recommendations

As the result of the discussions, workshop participants developed the following findings and recommendations.

Key Findings

- CI challenges: The explosion of large and diverse data sets is driving the decisions LFs make about technologies and software solutions. At the same time, CI is ever-changing and increasing in complexity, and making decisions about and adopting new solutions is complex within an operational LF environment.
- CI challenges: Although scientific networks have improved tremendously, some LFs still struggle to transfer data and access services, especially in remote environments.
- CI challenges: Integration, interoperability, and reuse of cyberinfrastructure solutions could be much improved. There is a natural tension between the community of CI software developers who seek commoditization where applicable and LF missions, which can be highly tailored. Thus, assistance is needed in promoting interoperability via trusted intermediaries.
- CI challenges: There is a natural tension between operations and maintenance needs and those of the end-user, both of which keep increasing over time.
- Collaboration: There are many opportunities to enable collaborations among different LFs or between LFs and CI projects. There is a set of common CI services that could be directly shared or leveraged across facilities. In some cases, lessons learned, best practices, architectures, and technical stacks and their configurations could be shared. In addition, CI should strive to move up the value chain toward more advanced services.
- Collaboration: By fostering collaborative and community efforts, LFs can potentially achieve economies of scale from working together on designing and deploying CI solutions, which in turn may lead to continuity for the solutions used and the ability to deliver a capability beyond the resources of a single LF.
- Collaboration: Fixed budgets for already committed resources and often incompatible timescales are significant barriers to collaboration between LFs and with other CI projects.



- Collaboration: Risk management and mitigation is an integral operational component of an LF that has a direct impact on research outcomes. Additionally, there are still barriers for adopting well-known solutions developed by potential competitors.
- Operational Practices: Management of CI facilities is hard because of mismatches between domain approaches to management and operations in a necessarily interdisciplinary environment (science and computer/engineering).
- Workforce: Hiring, retention, and advancement are particularly challenging at CI facilities because of funding uncertainty, differences between research and industry skillsets and work environments, and a lack of clear promotion paths. However, there are promising activities such as training and professional internships that could be leveraged.
- Workforce: Since hiring from computer science degree programs has not proven optimal, looking at candidates from other backgrounds may be useful.

Key Recommendation Actions:

The workshop participants recommended a number of actions to address the challenges faced by LFs. The recommendations below could be enacted through a combination of community efforts, facility peer interactions, and facility-CI project/platform expertise exchange as well as trusted entities such as dedicated centers.

- Mechanisms for CI discovery and opportunities for sharing of existing solutions, services, and training resources amongst the LFs and CI projects need to be supported.
- A common repository of knowledge about CI best practices, system descriptions, architectures, use cases, and core system tools should be created and made available to the broad community.
- Trusted intermediaries that can help navigate the ever-changing CI landscape, especially when migrating to new solutions need to be funded. Such entities can also assist in science-driven blueprinting of LFs before CI work begins.
- Communication, collaboration, and community-building efforts across LFs and CI projects should be fostered, and the benefits of belonging to a larger society of professionals need to be communicated.
- Research into new and effective methods for incentivizing collaboration and engagement with multiple facilities on joint projects needs to be supported.
- An effort to capture, analyze, and disseminate LF best practices and management techniques and to identify productive engagement opportunities between science, engineering, innovation, leadership, compliance, and other roles needs to be supported.
- Activities geared toward providing affordable training opportunities, helping structure career paths across facilities, recruiting talent starting with the undergraduate level, creating networking opportunities for technical (especially earlier-career) CI staff across facilities, and helping the community to understand the nature of CI as a career need to be supported.



Introduction

Overview and Goals

In 2015, the National Science Foundation (NSF) began supporting biennial workshops focused on cyberinfrastructure (CI) for large facilities (LFs). The workshops bring together large facilities and CI projects to share common experiences and challenges, discuss potential collaborations, and identify opportunities for leveraging CI within the two communities.

The 2017 CI for Large Facilities workshop found that “the need for, and benefits of, close interactions, collaborations, and sharing among the facilities and with the CI communities are well recognized, including the sharing of CI related expertise, technical solutions, best practices, and innovations across NSF large facilities as well as research facilities outside NSF (DOE, NIH, NASA, etc.).” Among the recommendations from that workshop was one to “foster the creation of a facilities’ CI community and establish mechanisms and resources to enable the community to interact, collaborate, and share” [2017 LF CI report].

The 2019 NSF Workshop on Connecting LF and CI aimed to continue and advance previous discussions, enable the exchange of CI solutions and challenges, and foster CI community building around NSF large facilities. It provided a forum to share ideas and experiences and to prepare for future CI research, development, and deployment that supports cutting-edge science. The major theme of the 2019 workshop was “Connecting Large Facilities, Connecting CI, Connecting People,” which emphasized the need to facilitate collaborations among LFs and CI projects and recognized the importance of the CI workforce to LF science missions.

The workshop was publicized at the 2019 Large Facilities Workshop held in April [2019 LF workshop], which included a day dedicated to CI: “Envisioning the Future of Facility Science & Cyberinfrastructure.” Participants at the event received a short survey about topics they would be interested in discussing at the September meeting. Based on the survey responses and discussions within the steering committee, a set of initial, high-level workshop goals were identified.

Specific goals of the 2019 workshop included:

- Identify CI challenges that facilities face when supporting their science missions;
- Explore opportunities and obstacles to collaboration between LFs;
- Examine non-technical challenges that influence CI development; and
- Develop ideas for enhancing the CI workforce and building a community of CI professionals.

Workshop Organization and Attendees

Steering Committee

The workshop was organized by a steering committee composed of CI experts from LFs and selected CI projects:

- Brian Bockelman, Morgridge Institute and HTCondor project [htcondor]
- Adam Bolton, National Optical Astronomy Observatory [noao]
- Tom Cheatham, University of Utah and Campus Research Computing Consortium [carcc]
- Ewa Deelman (PI and Chair), University of Southern California and CI CoE Pilot [cicoe-pilot]
- Tom Gulbransen, Battelle and NEON [neon]
- Kate Keahey, Argonne National Laboratory and Chameleon [chameleon]
- Marina Kogan, University of Utah



- Dan Stanzione, Texas Advanced Computing Center and the Leadership-Class Computing Facility [lccf]
- Daryl Swensen, Oregon State University and Regional Class Research Vessel [rcrv]

The 94 workshop participants included representatives from the NSF LFs, major NSF-funded CI projects, selected cloud providers, and NSF staff. Every effort was made to have at least one representative from each LF [lf_list, Appendix C]. For facilities with a distributed collaboration, representatives for the various LF platforms were invited. This was the case for the National Hazards Engineering Research Infrastructure [nheri], which includes a number of experimental facilities as well as different class vessels that make up the Academic Research Fleet [arf]. The list of attendees is found in Appendix A.

Workshop Structure and Activities

The major topics of the workshop built on discussions, findings, and recommendations from previous workshops and meetings. The agenda for the workshop was defined by the steering committee with community input. A spring 2019 survey advertised to LF workshop participants that attended CI day helped define the overall goals of the workshop by soliciting ideas for discussion topics and suggestions about people to invite to the workshop. A second survey sent to workshop participants prior to the meeting helped define particular panels and breakout sessions. The participants were also offered the opportunity to give a short “lightning” talk about their work, which resulted in 12 presentations. Finally, the workshop included a guided activity in which participants were divided into small groups and asked to answer two questions: 1) What are the most significant challenges faced by LFs or projects? 2) What are the most important problems a Cyberinfrastructure Center of Excellence could solve? The groups discussed the questions and were given the opportunity to share their answers verbally with all groups and to contribute their answers to the workshop notes.



*Kate Keahey (ANL) and Rafael Ferreira da Silva (USC)
lead the Guided Activity Session.*

Meeting participants were given the opportunity to provide input on workshop notes taken using Google Docs. Selected participants (Appendix A) were also asked to take notes during all the sessions to ensure that the discussions were captured. The raw, read-only notes are linked to the workshop agenda page (<https://facilitiesci.github.io/2019/agenda.html>).

To obtain a quick pulse of the community, workshop participants were asked to send a survey to their teams to gain information about knowledge sharing, incentives to following a CI career path, and feelings of belonging to the broader professional CI community. The analysis of the results is discussed in the ‘Participant and Practitioner Surveys Summary’ section starting on page 19.

The workshop also started a community-building effort by requesting participants to provide “CI calling cards” that included their professional information, pictures, and responses to questions about their recent CI frustrations and successes, both technical and non-technical. These calling cards greatly facilitated group and one-on-one discussions and are included in Appendix D.



This report is organized around the main topics of the workshop rather than the program components. Findings are formatted in *italics*, and recommendations are underlined. All workshop materials, including presentations, notes, photographs, and calling cards, are available on the workshop website: <http://facilitiesci.org>.

Theme I: Identifying cyberinfrastructure challenges that facilities face

Prior CI for LFs workshops identified the technical challenges of keeping up with acquisition; processing, and delivery of vast, heterogeneous, and dynamic data; tracking and managing technological changes; employing automation to streamline data processing and operations; developing well-informed and robust CI; and sustaining CI solutions over time. Although the various challenges were discussed throughout the meeting, the breakout session “What are the CI challenges that need to be addressed in the next 5 years to support LFs?” was dedicated to these discussions. In this theme, we highlight some of the challenges discussed in greater detail at the workshop.

The explosion of large and diverse data sets is driving the decisions LFs make about technologies and software solutions. The growing complexity of sensors and instrumentation, longer lifespans of LFs, and technological advances that make it possible to cheaply acquire and store data is leading to an overwhelming explosion in the volumes, acquisition rates, and heterogeneity of data. The data explosion not only affects short-term storage and long-term archiving but also decisions LFs make about technologies and software solutions. Many LFs must balance their budgets with priorities of long-term archival and real-time delivery of instrument and pre-processed data. For example, LFs that rely on relatively slow and expensive satellite uplinks to move data from instruments to data processing or archiving centers, such as IceCube and the Academic Research Fleet, must balance uploading high priority science data with quality of life issues such as internet access for the staff at a facility or on a ship.



Daryl Swensen, Oregon State University, reporting on the breakout session “What are the CI challenges that need to be addressed in the next 5 years to support LFs science missions?”

Although scientific networks have improved tremendously, some LFs still struggle to transfer data and access services. Many LFs feel that there are times when networking can still be challenging in modern, heavily connected environments. Problems can occur when performing high-bandwidth transfers between data centers across the continent or between online storage and backup facilities. An interesting discussion point was that bandwidth is not the only concern for the more sporadically connected facilities. There is also a need to “always be connected” to use online collaborative tools such as Dropbox, Google Docs, Microsoft One Drive, or when dealing with software licenses such as the Adobe software suite that requires an online network connection.

CI is ever-changing and increasing in complexity, and making decisions about and adopting new solutions is complex within an operational LF environment. Software development is also an area of concern for many LFs. Because many LFs have longer lifespans, the challenges most often brought up by participants related to long-term stability, supporting legacy code bases, and allowing



scientists to use a never-ending, constantly changing set of modern software tools. For LFs to operate smoothly, they must perform thorough requirement analysis while managing expectations from a funding agency and the community they serve. However, this requirement analysis often needs to be revisited many times during the lifetime of the LF. Several attendees also mentioned that compound interest on technical debt is expensive, and keeping and supporting older codebases is usually not a good option. Another software-related consideration is whether to use academic-grade software or commercial solutions, both of which could eventually bring prohibitive cost. The former solution might go away or become unsupported while the latter might lead to vendor lock-in and has its own sustainability challenges. Additionally, many LFs stated a desire to increase their use of machine learning algorithms for recognizing complex data patterns. However, state-of-the-art machine learning is available primarily in modern, community-driven software tool chains that present transition barriers to LFs.

Integration, interoperability, and reuse of cyberinfrastructure solutions could be much improved. One challenge discussed at the workshop was that facilities typically address CI challenges independently of each other and develop custom solutions in an uncoordinated manner. The result is that they are re-inventing similar solutions and miss opportunities to leverage work and knowledge. LF lifespan and different project timelines can make it difficult to identify and commit to long-term CI partnerships.

There is tension between operations and maintenance as well as the needs of the end user, the latter increasing over time. LF mission focus and the need to mitigate long-term risk often result in the desire to invest in in-house software—resulting in a plethora of bespoke systems—rather than adopting existing solutions or working collaboratively across LFs and large CI projects. Timelines and budgets under which LFs operate also may drive them to pursue expedient solutions rather than step back and examine potential new technologies and solutions. The same pressures may limit LF willingness to seek and maintain collaborations that can potentially increase CI reuse and new CI adoption.



Patrick Brady, University of Wisconsin Milwaukee, gives a lightning talk about “SCiMMA: Scalable Cyberinfrastructure to support Multimessenger Astrophysics.”



Recommendations

The workshop found that for LFs to develop, adopt, and maintain robust CI requires planning that is flexible enough to allow agility in moving forward in incremental steps over the LF lifespan. Planning, including quantifying the potential benefits and the cost of making changes, is important because of the difficulty in making changes to Major Research Equipment and Facilities Construction projects (MREFCs) in progress even when the technology landscape is continuously and rapidly changing. During the conceptualization phase, it is especially important to bring in CI expertise, people with a frame of reference for operations, and people with a broad overview of current robust CI capabilities. During the LF operational phase, it is critical to re-integrate with the broader CI community to leverage each other's strengths and develop approaches to change management. There is a need to create opportunities for CI discovery and sharing of existing solutions, services, training resources amongst the LFs as well as CI projects.

A somewhat obvious, yet important, recommendation is to facilitate the use of best practices and effective processes. Creating a common repository of knowledge about CI best practices, system descriptions, architectures, use cases, and core system tools—e.g., a taxonomy of logical architectures for LFs—is desirable. The CI community should work towards a set of guidelines on creating well-developed documents that capture architectural details, rationale, and alternatives studied before reaching the decision on CI. The knowledge repository could also include information about common software and services, data formats and ontologies, data curation and preservation techniques, and reproducible results. The ability to access community wisdom will become increasingly important as the principles of findable, accessible, interoperable, and reusable (FAIR) [fair] data are adopted by more scientific communities. The attendees were somewhat divided as to whether the knowledge base should also include evaluations and recommendations of some tools over others. Finally, a yearly, focused gathering of CI/LF senior architects may be a very effective way to complement (but not replace) the common repository as discussed above.

There was a lot of interest and recommendations for a trusted entity, collection of trusted entities, or Centers of Excellence. In general, the agreement was that the ever-changing and increasing complexity of CI raises the need for trusted intermediaries to help navigate the CI landscape. These intermediaries would assist in science-driven blueprinting of LFs before CI work begins, provide a mechanism to “matchmake”



*Margaret Johnson,
University of Illinois at
Urbana Champaign, gives a
lightning talk about
“Enabling Multi-Instrument
Pixel-Level Science with A
High Throughput
Computing, Data”*



between LFs and CI projects, and decide when the partnership works for both or could be applied across projects. A more complex recommendation was for the trusted entity to help clarify the distinctions, possibly case-by-case, between common/commodity CI solutions and those that are domain-specific and to also clarify the distinction between wants and needs for facilities.

Finally, commercial cloud solutions provide a number of capabilities (computing, archival storage, etc.) that can contribute to the LF science mission. However, the published cloud costs often appear cost prohibitive. The community could benefit from an effort to work with all cloud vendors to explore whether the commercial cloud can contribute to LF and CI in a cost-effective manner and provide the support that the science community needs. A potential area of exploration, which would potentially increase sustainability of facilities and researchers who create scientifically valuable data sets and algorithms, are new business models that incentivize cloud vendors to share their compute revenue with the creators of the content/algorithms subscribers use.

Theme II: Exploring the opportunities and obstacles to collaboration between LFs

Collaboration is essential to support and strengthen scientific progress. Bringing additional expertise and resources together can enable research endeavors that were not otherwise possible. On the other hand, when not properly managed, collaborative efforts may raise potential risks, challenges, and inefficiencies. Several discussions about opportunities and challenges to collaboration between LFs took place during the two-day workshop, including the panel discussion "Shared CI Services: Opportunities and Challenges." In this section, we summarize and highlight findings and recommendations for these opportunities and barriers to collaboration.

There are many opportunities to enable collaborations among different LFs or between LFs and CI projects. Although different facilities target different scientific problems from distinct science domains, there are a set of common CI services that could be directly shared or leveraged across facilities or about which lessons learned, best practices, architectures, and technical stacks and their configurations could be shared. Notably in the Participant and Practitioner Survey (see Appendix for full results), half of the respondents selected somewhat likely or very likely when asked about the likelihood of their facility/project being willing to contribute expertise to a shared services efforts. Potential common CI services include compute, storage, authentication, data discovery, data archive, disaster recovery preparedness, deployments, maintenance, and networks.

Current models for sharing resources among facilities include: (1) organizational (e.g., ITIL/COBIT model [itil] for shared IT service & governance), (2) mission-driven (e.g., DOE-funded CI centers), (3) vendor-oriented (e.g., commercial clouds), (4) partnership-based (e.g., partnership agreements with major NSF-funded CI centers), (5) collaborative (e.g., facilities may share CI with one another), and (6) ad hoc (e.g., people who occupy similar positions at different LFs informally communicate or meet at common conferences such as the forum AskCI [askc] to share recent experiences and opinions). The first three models are currently achieved via internal organization and economic incentives but have very limited scope. Partnership-based agreements with CI centers have demonstrated an impact on domains such as cybersecurity (Trusted CI [trustedci]) and science gateways (SGCI [sgci]).



Panel on Shared CI Services Opportunities and Challenges, Moderator: Adam Bolton (National Optical Astronomy Observatory).

Panelists: (from the left) Pamela Hill (National Center for Atmospheric Research), Von Welch (Indiana University and Trusted CI), JJ Kavelaars (National Research Council of Canada), and Michael Zentner (University of California San Diego and Science Gateways Community Institute).

By fostering collaborative and community efforts, LFs can potentially achieve economies of scale from working together on designing and deploying CI solutions, which in turn may lead to continuity for the solutions used and the ability to deliver a capability beyond the resources of a single LF. Collaboration can be achieved by sharing expertise, connections, and efficient processes that allow LFs to focus on their core competencies. For instance, NOAO [noao], LSST [lsst], and GEMINI [gemini] projects share several CI services, mostly focused on computing and data management, across their institutions. In addition to technology components, sharing competence, community, knowledge, and awareness also aggregates high value to the collaboration and may lead to improved project sustainability, funding acquisition, and workforce development. Furthermore, commonality in architectures in what are often niche academic spheres, can drive a tool or service in a direction that benefits all and develops a wider, more skilled community. Although this can occur for both commercial or open source tools and services, it is particularly likely to happen with open source resources for which LF CI staff are the primary developers. To establish such collaborative efforts among the diversity of scenarios for LFs and CI projects, several challenges need to be addressed, as discussed below.

Risk management and mitigation is an integral operational component of a LF that has a direct impact on research outcomes. Although outsourcing less specific and critical CI services (e.g., by leveraging other CI systems and platforms for computing and data management) may improve service robustness and availability, there is a perceived risk when relying on external software, hardware, and/or connectivity services. For instance, moving large data sets over the wide area network may be challenging due to low network connectivity and may also raise security issues. Accessing shared computing platforms requires predictable job execution and turnaround times as well as the ability to support complex model configurations (e.g., multiple job steps, in-situ analysis, data assimilation, machine learning components, different software environments, etc.).

There are still barriers for adopting well-known solutions developed by potential competitors. For instance, a common approach in CI software development is to redevelop solutions rather than tailoring community software to specific needs. This approach is typically motivated by the need to host solutions



developed in-house (software development path is customized to the LF need), lack of documentation, and/or long-term sustainability. The lack of formal common forums for collaboration and sharing of experience also inhibits competence sharing.

The timescales at which LFs operate is another hindrance to collaboration. It becomes incredibly complex to share information or collaborate with quickly changing technologies, especially when processes and budgetary structures do not support refactoring, evolution, sharing, interoperability, planning upgrades, and new deployments. Those factors are further complicated by the fact that simultaneous old and new deployment periods are inevitable during the upgrade of an operational facility.

Finally, there are both monetary and physical barriers that impose constraints on collaborations. CI practitioners working in remote environments such as ships or sites that are beyond connectivity simply cannot regularly communicate with others. Additionally, LFs have limited resources that are already allocated to serving their target users. Developing new relationships, forums, and other vehicles for collaboration is costly in the immediate timeframe, even if it will reap a multiplication in long-term rewards. Despite the above discussed willingness to share expertise, respondents to the Participant and Practitioner Survey were ambivalent when asked about the likelihood of their facility/project being willing to contribute resources (provide storage, compute services, personnel). And when asked what obstacles they foresaw in implementing a shared services effort, the largest portion of respondents (40.63%) indicated funding and personnel as the main obstacles.

Recommendations

Given the above, there is a clear opportunity and long-term value for the community and funding agencies to put in place mechanisms that enable collaboration. Excellence in collaboration requires trustworthy people, processes, and technology. Increased risks in collaborative projects can be mitigated by utilizing existing trusted relationships and working to actively develop new ones.

In the immediate term, there are some relatively simple activities that could be undertaken. For instance, sharing existing technical architectures in a form that makes them discoverable and reusable would allow LFs to recognize which other LFs might have relevant expertise to share. Additionally, it could inspire consideration of similar architectures while benefiting from lessons already learned. Similarly, the sharing of expertise regarding open source licenses, community data standards, and software development best practices were all voiced by the workshop participants as simple but tangible potential products of value if shared in a useful manner.

The results of the Participant and Practitioner Survey showed a willingness (within given caveats) to share expertise alongside the acknowledgement and calling out of the obstacles of limited funding and personnel time. Thus, it is plausible that tangible funding should be made available specifically to support collaborative sharing efforts. However, more research is needed to carry out a cost benefit analysis of sharing models before such a recommendation be made.

Lastly, as a longer-term effort, external support structures could foster and facilitate increased collaboration. For instance, external support could facilitate and foster possible community collaborations that appear to have the long-term potential to justify initialization cost. These collaborations will likely fall along a specific CI theme. This support could include maintenance of mailing lists, organization and hosting of monthly telecons, planning, recording, and distributing webinars that share LF architectures and lessons learned.

Other instantiation of external support could include a facilitator curating a special issue publication showcasing technical architectures (perhaps themed according to a stage of LF data lifecycle) or a publication surveying emerging technologies. Whether the latter would be written solely by external evaluators with the technical expertise to do so, or by LF staff could be domain dependent. Ultimately, these publications could serve as a source of collated and reviewed material that saves each LF from



carrying out its own in-depth evaluation of every new technology that appears. Clearly, these evaluations would need to clearly detail the use case context of the evaluation.

Finally, the workshop attendees recommended that a team dedicated to short-term projects involving multiple LFs (and requiring LFs to provide their own effort) would be highly effective. This final recommendation, however, is likely the most expensive.

Theme III: Examining non-technical challenges that influence CI development

Several non-technical issues and challenges that influence CI development for LFs and, more broadly, for large NSF CI projects, were identified during the discussions at the workshop. This section captures those discussions and recommendations from participants about how to overcome those challenges.

There was also a significant amount of discussion around models for a trusted entity or set of entities that could facilitate or nucleate community efforts, facility peer interactions, facility-CI project/platform expertise exchange, or provide services to the CI community. The example of the collaboration between the Cyberinfrastructure Center of Excellence Pilot [cicoe-pilot] and NEON [neon] was discussed as one possible type of a one-to-one engagement. The participants also shared ideas about other useful tasks that the community, LFs or other entities could perform to improve the LF CI development. Among them were:

- perform an assessment of all/many LFs to identify common needs and problems,
- evaluate new technologies and help LFs understand which are applicable to their use cases,
- provide training and recommendations while facilities stand up the hardware and software,
- maintain expertise in specialized areas (e.g., Internet of Things, workflows, data modeling, data archiving) where individual facilities cannot afford to do because of episodic rather than constant needs,
- provide opportunities for engagement between LFs and CI providers at key points in the operational timeline to meet needs for general consulting or for specific solutions/services,
- develop expertise and consulting around standards (not act as an auditing body though) e.g., evaluate NSF mandated proposal requirements with respect to sharing CI for LFs, and
- provide an ‘army’ of skills.



Breakout session on “What are the non-technical issues that influence CI development and how they can be collaboratively addressed?” The session was led by Susan Sons (Indiana University) and Douglas Thain (University of Notre Dame).



The management of CI facilities is difficult because of mismatches between domain approaches to management and operations in what is necessarily an interdisciplinary situation (science and computer/engineering/CI). Some of the challenges in operational practices are enumerated below.

There is a tension between adopting new technology/innovation/improvement and maintenance/feasibility/sustainability of CI. One of the operational challenges for large CI entities is to develop a coherent strategy for achieving balance between sustaining what has been developed vs. engagement in new activities with the consideration that resources might also be needed for urgent activities. There is often a lack of long-term planning for ensuring these different needs are kept balanced and for dealing with legacy infrastructure.

There is a lack of clear understanding about the ideal organizational structure within large CI facilities. The teams are often highly distributed and frequently exhibit “split agency” characteristics, meaning there is a mismatch between individual and organizational objectives resulting in conflicting demands and expectations. Sometimes, CI teams serve multiple PIs with conflicting needs and desires. Another concern is balancing the organizational structure to cater to domain science expertise and computing/CI expertise.

Staffing skills and knowledge transfer are a challenge faced by most facilities. Often, large CI facilities are heavily matrixed and are trying to maintain a small staff with many specialized skill sets. When turnover occurs, knowledge is lost. This is especially problematic when a developer leaves behind systems that were created using specialized knowledge and unique skills. How to manage these systems in a way that reduces loss of maintainability is a significant challenge. The lack of communication among technical staff and between technical and non-technical staff also impedes smooth knowledge transfer during successions. Insufficient tacit knowledge capture and lack of automation in knowledge capture are serious challenges for long-lived CI projects and LFs.

Recommendations

There needs to be an effort to analyze and distill best practices and management techniques, and for identifying productive engagement opportunities between science, engineering, innovation, leadership, compliance, and other roles that would be instrumental for LFs.

There needs to be support for research focused on successful staffing models (science vs. computer/engineering/CI) and resource management practices, and a way to share those outcomes with NSF, LFs, and large CI projects and organizations. It would be beneficial to create a body of best practices around well-integrated CI engineering organizations within science organizations.

There is a need for creating trusted entities that can follow relevant technology trends, evaluate new solutions in the context of LF CI, make recommendations, and provide training.

Theme IV: Developing ideas for enhancing the CI workforce and building a community of CI professionals

Discussions about enhancing the CI workforce and building a community around CI permeated most workshop sessions. Developing and growing a community of CI practitioners available to support the science missions of LFs was identified as one of the major challenges faced by workshop participants. The word *community* can be difficult to define, particularly with no unifying entity to establish an identity for the community or to govern its actions. Workshop participants agreed that a community is a group of people who share a common language and common function or purpose, but the boundaries are not always clear. For example, a university, LF, or organization like XSEDE [xsede] could be a community,



but so can a department within a university, a professional society spanning multiple organizations, or a group of people with similar interests from a variety of different locales.

Communities should be venues to share expertise and experience, despite the challenges in drawing strict lines around communities. They should have a centralized way to communicate and exchange ideas. People within communities should see mutual benefit in belonging and should be self-selecting. A Cyberinfrastructure Center of Excellence could help people self-organize into the much-needed community of CI practitioners.

Today's CI practitioners may have found their positions at LFs via an internship program (e.g., NCSA Cyberinfrastructure Professional Intern Program [cip, mate]) or a contract position where they were assessed for fit before being offered a permanent position. Of those who receive a permanent posting, workshop participants estimate about two out of three stay.

Workshop participants found many challenges associated with recruiting:

- Compared to industry, some LFs have requirements or working conditions that may make them less appealing. For example, working at the South Pole requires tolerance to cold and isolation, and working on research vessels may exclude candidates who suffer from chronic medical conditions if they need continual specialized care. Some LFs require employees to work remotely for long periods of time, away from home and family.
- Typical computer science curriculums may not adequately prepare students for work in an LF.
- Graduating students may not even be aware of LFs. Campus recruiting fairs may not think to include LF representation, and career services' counselors may not be aware of LF opportunities.
- Job descriptions may not accurately represent the role because the institution's human resources may require that postings fit titles and templates that suit their needs, rather than the needs of the hiring manager. For example, some participants said they had to name positions with more run-of-the-mill titles, such as software developer, when the role is substantially more nuanced than that. This can lead to misunderstanding about the requirements of the job and frustrate the hiring process.
- Because LFs tend to operate differently from their parent institution/university, hiring managers may have difficulty in explaining or justifying hiring, compensation, or promotion needs to a human resources department that primarily deals with standard university hiring and retention.
- Recruiting international employees may prove difficult or impossible due to the institution's export control practices or reluctance to, or difficulty with, managing visas.
- Salaries are not competitive with those offered in industry. Google, Microsoft, and other high-profile tech companies can pay far more than an NSF budget allows.
- Funding uncertainties such as limited term cooperative agreements make it hard to hire and retain quality CI professionals.

Retaining staff is also a challenge due to a number of factors:

- Career paths are not well defined. There is no common set of job titles and descriptions, paths for attaining raises and promotions, or sets of necessary certifications or qualifications. Instead, people tend to "fall into" CI roles, which makes it harder to justify positions or the value of the individual employee (and subsequently his/her raise or promotion) at performance review time. Because the CI practitioner role(s) has not been formalized across the market, institutions/universities are hampered in providing advancement, development, and proper compensation.
- CI practitioner roles tend to be filled by individuals of the same demographic. Few women and people of color appear in these roles. Those who do may be less likely to stay if they feel isolated from or socially different from their co-workers or perceive imbalances based on race and sex.
- Many teams at LFs are geographically distributed. They may work at different time zones, speak different languages, and communicate less frequently or via less rich means (i.e., not face-to-face).



Although this diversity has many benefits, it also proves challenging for coordination, communication, and establishing a sense of shared identity and community among team members.

- Some LFs require employees to travel for extended periods of time, and science goals may need to be addressed on accelerated schedules. Both can make it difficult to carve out time for employee development and training.
- Sometimes there is a disconnect between the scientists and CI practitioners. They may not share the same language and will necessarily approach problems from different perspectives. This may make common understanding, and consequently achievement of goals, challenging. To avoid frustration, time should be taken to either cross-train, find individuals with a blend of these skills, or foster communication and understanding.

Although there are many challenges around this workforce issue, there are many things that work well and can be further exploited to improve the situation. For example, hiring managers can appeal to work/life balance, which workshop participants felt was more favorable than in the private sector, or to the science mission which, to some candidates, would make working at an LF meaningful and enriching. Many training opportunities already exist, including Open Science Grid's user school [osg], TrustedCI training, and the Linux Cluster Institute. Some networking opportunities exist as well, including the workshop that this report discusses, the Campus Research Computing Consortium (CaRCC) [carcc], and the Coalition for Academic Scientific Computing (CASC) [casc]. Additionally, LFs such as the Academic Research Fleet [arf] have a successful MATE program [mate].

Additionally, one could create a trusted entity or set of entities to help:

- Research and provide a repository of information and resources around developing community and the CI practitioner workforce, including:
 - mailing lists around topics of interest,
 - lists of sub-communities and their annual meetings,
 - best practices, guides, and templates on matters of community building, hiring practices, employee development, etc., and
 - hosting a CI job board.
- Serve as a connector to bring people together by listening to and knowing the larger LF community so that members can make meaningful connections with and receive help and guidance from peers. This could also involve serving as a "birds of a feather" space where people can congregate virtually or in person and self-organize into subcommunities.
- Create opportunities for engagement and knowledge sharing, such as hosting a version of StackExchange for CI practitioners; facilitating or helping establish mentoring programs; offering or collaborating on the development of training on all aspects of CI; facilitating or helping establish an exchange program where CI practitioners work at other LFs for a time; offering other networking and knowledge-sharing events.
- Work with universities to make students and career services more aware of job opportunities in LFs.

Recommendations

There needs to be an effort to help facilities meet workforce requirements by providing affordable training, creating career paths across facilities, recruiting talent starting at the undergraduate level, and helping the community understand the nature of CI as a career. One recommendation is to focus on hiring domain science graduates to work in computing/CI areas. It is also important to recognize that research CI professionals often remain in their jobs even if the pay is not competitive because they enjoy the work and science. It would be beneficial to create an “academy” for improving computing/CI skills that are broadly applicable for LFs.



A trusted entity could also provide staff training on emerging solutions that a specific LF does not have time to provide. Articulating career paths for CI professionals is also important. This could be done by interfacing with undergraduates to make career paths known and by talking to career services at universities. It could also be beneficial to research factors that cause people to stay or leave CI jobs, potentially designing and analyzing exit surveys. Exploring avenues like cross-facility job postings or offering CI experts the opportunity to work at other facilities for a period rather than leaving for industry, could also help. It may be beneficial to develop an internship funnel for under-represented communities and facilitate Research Experiences for Undergraduates (REU) programs related to CI.



*Panel on Workforce Development and Retention,
From the left: Sharon Broude Geva (University of Michigan); Tom Cheatham (University of Utah, panel moderator), Frank Wuerthwein (University of California San Diego and CMS), Rachel Adams (UC Boulder), and Ji*

Cyberinfrastructure Calling Card Summary

Prior to the workshop, participants were asked to complete a CI calling card that was used to help participants meet each other virtually before the workshop and to facilitate conversations about topics of mutual interest between participants during the workshop. The calling cards asked participants to briefly comment on three topics:

- a CI accomplishment they achieved over the past year,
- a CI frustration or challenge they faced over the past year, and
- a non-technical CI issue or success they would like to share with workshop participants.

CI accomplishment	Release of the IceProd 2 workload management system. This release simplifies using it to manage experimental data processing as well as simulation and is a step toward becoming a user facing platform.
CI frustration or challenges	Adapting a 15 year old code base to fit into contemporary CI systems and workflows.
Non-technical CI issue or success	Finding the time and resource to experiment with new computing models and analysis techniques.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure


Steve Barnett
barnet@icecube.wisc.edu
University of Wisconsin-Madison


IceCube Neutrino Observatory
<https://www.icecube.wisc.edu>

An example of a CI Calling Card.



Participants submitted 61 calling cards, which were analyzed after the workshop for common themes or patterns. Calling card comments were grouped into two main categories: (a) accomplishments and successes and (b) frustrations, challenges, or issues. These were further subdivided into themes that manifested during the analysis process.

Themes within comments about accomplishments and successes:

- *Deployment of something new (e.g., systems, services, infrastructure).* Examples included deploying clusters around the globe or on research vessels and creating system assessment tools and metrics that allow reporting to stakeholders.
- *Improvement, expansion of an existing system, service, or infrastructure.* Examples included expanding storage, automating cyber-physical experiments, and provisioning 10 Gbps fiber to telescopes in remote locations.
- *Continuing to do good work throughout the year.* This was notable because often day-to-day activities may not be characterized as achievements due a definition of accomplishments as novel and set apart from the day-to-day. However, duties done well are vital to the successful operation of any type of organization, including large facilities. Examples included ensuring data quality is assessed and monitored; finishing construction and launching on time and within budget; and balancing the needs of stakeholders while addressing challenges.
- *Communication-oriented activities,* specifically around bringing people together, were mentioned frequently. Examples included helping distributed units of a large facility connect and feel more like a network, partnering with peer organizations to make improvements, participating in related communities (e.g., EarthCube [earthcube], ESIP [esip]) to advance shared goals, building effective teams, and facilitating collaboration between researchers and developers.

There were other types of successes mentioned that did not fall into any of the above themes, such as instituting an efficient procurement process, meeting the needs of funding partners, successfully meeting a standards base (e.g. Open Data Act, ITAR) or applying one, and leading development efforts.

Themes around frustrations, challenges, or issues:

- *Frustrations around data* included those around sharing, disseminating, and accessing data; data management and curation practices; and the impact that increasing data acquisition rates have on things like storage and management.
- *Budget woes*, which include scaling up without the budget spinning out of control as well as mismatches between operations timelines and funding cycles that may negatively impact the ability to respond to evolving technologies.
- *Technical issues*, which can be further divided into:
 - *Planning, systems analysis, and design.* For example, learning there is more to manage (e.g., support, migrate, refactor) or more is needed to satisfy needs than was originally anticipated or planned for.
 - *Old vs. new technology*, such as responding to evolving needs or keeping older systems up-to-date with current technology.
 - *Heterogeneous technology*, such as integrating different technologies based on differing needs; the need for specialized or customized tools because one size does not fit all; or working with researchers that use different platforms, operating systems, software, etc.
 - *Shared services and distributed services*, such as the need to gain the attention of shared resource providers; or managing the details of distributed storage allocations.
- *Communication-oriented challenges*, including changing culture and mindsets around a variety of things (e.g., motivating scientists to publish data, encouraging operational thinking, redefining



value); contending with conflicting or changing needs of stakeholders; and communicating with administration on mission focus or technology costs.

- *Workforce-related challenges* include a sense of disconnect between CI professionals and scientists or needing to blend the two skill sets; recruitment and retention of professionals in a competitive technical market; or overtaxed staff due to multiple funding obligations and reporting lines.

Other challenges, frustrations, or issues mentioned that did not fall into one of the above themes included desires for more information on best practices; limited bandwidth and reliability of internet connectivity on research vessels; and complying with multiple disparate standards.


Participant and Practitioner Surveys Summary

Workshop Participant Survey

Thirty-two workshop participants responded (43% response rate, since we did not send the survey to NSF staff) to a pre-workshop survey that focused on opportunities for collaboration and sharing expertise across LFs.

Respondents indicated LF-shared services in which they would consider participating. The survey provided the following options (multiple choice): Identity management, Storage-as-a-Service, Disaster recovery, Computing-as-a-Service, Data discovery, and the ability to write in additional services. The graph below shows that many participants saw the benefits of sharing these services.

What shared services across Large Facilities would you consider participating in (if any)?




Respondents also wrote-in additional services they would like to potentially participate in. These “write-ins” included: shared personnel services to support contracting, training and workforce development, diversity, organizational and policy aspects, satellite communications, global integration of computing and data, global content delivery system, and resource scheduling.

Respondents were asked about the likelihood of their facility/project being willing to contribute resources (provide storage, compute services, personnel) to the shared efforts. For the response, they were offered a 5-point Likert scale (from 1: very unlikely to 5: very likely). For the willingness to contribute resources, the response mean was 2.94 and median 3, indicating that respondents are rather ambivalent about committing resources to a shared service.



When asked which resources their facility/project would be willing to contribute to these shared services, most respondents did not answer the question or indicated the inability to share any/many resources. Only eleven respondents indicated a few salient types of resources their organizations would be willing to share:

If any, what resources would your facility/project be willing to contribute to these shared services efforts?




On the other hand, when asked about the likelihood of their facility/project being willing to contribute expertise to the efforts of shared services, the respondents were much more optimistic, with mean of 3.69 and median of 4 on the 5-point Likert scale. Overall, respondents were much more willing to contribute their expertise to a shared service rather than resources. The high median indicates that half of the respondents selected somewhat likely or very likely, signaling a high level of buy-in. This is well-illustrated by this participant quote: “I would think an expert-exchange program would be excellent – if nothing else a forum for knowledge exchange.”

When asked what specific expertise their facility/project would be willing to share, the respondents indicated many diverse types of expertise, including identity management, network engineering, cybersecurity, storage, statistics, data science, high-performance computing, remote, shared CI, standards, restful service layers, user management, cloud computing, replicability of experiments, user access, and personnel.

When asked what obstacles they foresaw in implementing a shared services effort, the largest portion of respondents indicated funding as the main obstacle, at 29%. The other salient categories are shown in the graph below.

What obstacles do you foresee in implementing a shared services effort?



One respondent stressed the lack of willingness to collaborate as the main obstacle and pointed out its relationship to the life cycles of LFs: “None other than the willingness to collaborate. In general, large



facilities projects are very insular in their initial phase because they are totally inward focused on getting the job done and perceive everything shared as a risk. So, it is a matter of control, and risk mitigation. As they mature, they tend to be more likely to collaborate. Except for the fear of ‘their money’ being spent on a collaborative effort that they do not control. So, a fear of losing funding and control are the real obstacles.”

A large proportion of participants indicated online platforms like discussion forums and chatrooms (Slack, Microsoft Teams, etc.) would be helpful for discussing CI-related issues (28.13%). One respondent illustrated the hopes for such an online space quite well: “single point (website?) for discoverability of current CI providers, initiatives, and solutions.” Other popular options included continuing annual CI meetings (6.25%) and a combination of in-person meetings and digital platforms (9.38%). Another participant indicated that the platform is less important than the motivation to contribute: “Any platform will work, but ‘helpful’ will happen only if people use the platform. Infrastructure and shared service discussions require domain expertise and IT expertise and topical abstraction up and down, requiring regular and repeated discussion at many levels and probably multiple platforms. How do we motivate people to engage in the conversation?”

Next, respondents were asked about shared workforce development activities across NSF LFs. They were first asked what aspects such shared efforts would need to have for participants to consider participating in them. Funding or attractive costs were again the most prominent category (12.5%). Other salient aspects of the shared workforce development effort included activities that enhance the ability for LF employees to share knowledge, and HR topics/retention for supervisors.

When asked how likely their facility/project would be willing to contribute resources to a shared workforce development effort, respondents did not feel strongly (mean=3.28, median=3 on the 5-point Likert scale). However, there were more highly positive responses than in the question about a shared service.

When asked what resources their facilities/projects would be willing to contribute to a shared workforce development effort, most respondents did not answer or indicated their facilities did not have relevant resources to share (50%). Those who specified concrete resources mentioned human resource/developers/personnel (9.38%) and expert/supervisor time (6.25%). There were also many unique responses offering both specialized resources and expertise.


Again, respondents were more willing to contribute expertise than resources to a shared workforce development effort, with a mean of 3.69 and median of 3.5 on the 5-point Likert scale. When asked what expertise they would be willing to share, the vast majority (59.38%) left the field blank or responded with “unsure” or “none.” The two salient types of expertise mentioned by other respondents were broad experience in CI development and leadership (12.5%) and domain expertise and experience working with remote, shared CI (6.25%).

Finally, the respondents indicated that the most likely obstacles in implementing a shared workforce development effort would be similar to the obstacles perceived for implementing the shared services, most prominently the lack of funding, time, or personnel (47%). Other obstacles mentioned included lack of shared skills, divergent goals, competition with industry, scheduling, agreement to collaborate, and connectivity.

The workshop organizers also ask the participants if they had advice for hiring CI practitioners. Twenty-four participants answered this question in a free form paragraph. The graph below shows an analysis of the results.



What advice do you have for someone who is trying to hire a cyberinfrastructure developer or administrator?



The advice of training people by starting with junior people or people with a non-CS background was most frequently given. The responders also indicated that it is important to set a realistic expectation of what needs to be accomplished and what the personnel can expect. Engaging in the science mission and providing an exciting and flexible work environment were also indicated, although at a lower level, with 3, 2, and 1 response each. One participant also acknowledged that staff turnovers will happen and that you need to “hire personnel with strong organizational and documentation skills, and ability to think in multiple contexts.”

Cyberinfrastructure Practitioner Survey

The workshop also conducted a very brief survey of CI practitioners associated with the workshop participants, with the aim to better understand knowledge sharing in the community. Seventy-eight responses were received with over 15 different LFs participating. Most of the responses came from CMS (14), the Academic Research Fleet (11), and NHERI (5) LFs. Six members of the Open Science Grid, which is a distributed CI serving LFs such as ATLAS, CMS, IceCube, and LIGO, also provided input.

The workshop organizers also asked the CI practitioners, “Do you feel like you are part of a community? If so, what is it and how do you participate in your community?” Sixty-eight respondents answered yes, and 10 answered no—that they did not feel part of the community. Not surprisingly, members of the CMS LF, who have a large number of collaborators and who were the major responders to the survey felt a great sense of belonging to a greater community. An example of a response from that community included “*Yes. The USCMS program creates a strong community of computing professionals and maintains good internal communication.*”

Based on the responses to the second question, we aimed to distinguish between people belonging to a particular project, LF, rather than the broader CI community, for example, by attending conferences that are organized by outside entities. Although this is not a conclusive result because of the open-ended nature of the question, we would estimate that over 42% of respondents felt they belonged to a broader community of CI professionals.

Since much of the workforce issues revolved around recruitment and retention, we wanted to explore why the CI practitioners pursue a career in academia or research laboratories. We asked, “What incentivizes you to work in cyberinfrastructure for science and research as opposed to industry?” The respondents were able to provide free text answers, and we categorized them as follows: having a sense of purpose, helping science and the world, the ability to conduct research and solve complex problems, having intellectual




freedom to pursue new ideas or contribute to open source projects, and working in an academic environment or having a flexible work culture or job stability. In some cases, CI practitioners were domain scientists that develop CI to meet their science objectives. An example of such a response was: "I am a professor of physics and care about the domain science deeply - not just providing general cyber-infrastructure."

In some cases, the responders indicated multiple incentives for working in CI. The graph below summarizes the results. It is clear that having a sense of purpose and helping science was the foremost incentive for the responders (30). One respondent answered the question with "Belief that what we are doing helps the world." Being able to conduct research and solve complex problems was the second most cited reason (21): "Simply put, it isn't boring. Science disciplines evolve and the needs and demands for infrastructure evolve with them. Matching technology to those needs is constantly changing and in need of "engineering" effort, so there is plenty to do!"


Although the sample of people surveyed was limited, the results do indicate that working in CI can provide a meaningful and attractive work environment. This finding can also potentially help recruit new personnel to careers in CI.

What incentivizes you to work in cyberinfrastructure for science and research as opposed to industry?"



We asked CI practitioners about where they get help when they run into problems and issues. The respondents indicated a variety of sources (they were asked to check all the categories that applied):

When you run into problems and issues, where do you get help today?




Over 90% of respondents turn to Google Search for help with the CI issues, and over 70% make use of the general internet forums. On the other hand, the fact that 82% of respondents turn to CI colleagues at their institutions (and 70% to those at other institutions) suggests that CI-specific resources would be in high demand.



We find similar trends with where the CI practitioners share CI best practices they developed or come across:

Where do you share the best practices of CI that you have developed or come across?



Clearly, providing a centralized CI-related forum for both searching for and sharing the best CI practices would be very beneficial to the community.

Finally, the respondents were asked to specify their favorite CI resource. There were 32 different answers in total. The most common answers were Google and GitHub, which were both mentioned by 8 participants (10.5% of total participants). OSG was mentioned 5 times (6.6%). Stackoverflow and Campus Champions were both mentioned twice (2.6%). The remaining answers, very community specific, were all only mentioned once (1.3%), like CERN, CaRRC, RVTEC, and so on.



References

- [2017 LF CI report] 2017 NSF Large Facilities Cyberinfrastructure Workshop Report
<https://facilitiesci.github.io/assets/reports/facilitiesci-workshop-report-11-17.pdf>
- [2019 LF workshop] Pamphlet about the 2019 Large Facilities Workshop
<https://www.largefacilitiesworkshop.com/wp-content/uploads/2019/03/19LFWpamphlet.pdf>
- [arf] Academic Research Fleet <https://www.unols.org/>
- [askc] Ask Cyberinfrastructure <https://ask.cyberinfrastructure.org/>
- [carcc] Campus Research Computing Consortium <https://carcc.org/>
- [casc] Coalition for Academic Scientific Computation <https://casc.org>
- [chameleon] Chameleon <https://www.chameleoncloud.org/>
- [ciccoe-pilot] Cyberinfrastructure Center of Excellence Pilot <https://ciccoe-pilot.org/>
- [cip] Cyberinfrastructure Professional Intern Program
- [earthcube] EarthCube <https://earthcube.org/>
- [esip] Earth Science Information Partners <https://www.esipfed.org/>
- [fair] Wilkinson, Mark D., Michel Dumontier, IJsbrand Jan Aalbersberg, Gabrielle Appleton, Myles Axton, Arie Baak, Niklas Blomberg et al. "The FAIR Guiding Principles for scientific data management and stewardship." *Scientific data* 3 (2016).
- [gemini] Gemini Observatory <http://www.gemini.edu/>
- [htcondor] HTCondor <https://research.cs.wisc.edu/htcondor/>
- [itil] Robert R. Moeller, Executive's Guide to IT Governance: Improving Systems Processes with Service Management, COBIT, and ITIL, 27 February 2013, ISBN:9781118138618, Online ISBN:9781118540176, DOI:10.1002/9781118540176.
- [lccf] Leadership-Class Computing Facility <https://lccf.tacc.utexas.edu/>
- [lf_list] <https://www.nsf.gov/bfa/lfo/docs/major-facilities-list.pdf> as of November 05, 2019
- [lsst] The Large Synoptic Survey Telescope <https://www.lsst.org/lsst>
- [mate] Marine Advanced Technology Education <http://www.marinetech.org/internships/>
- [neon] The National Ecological Observatory Network <https://www.neonscience.org/>
- [nheri] National Hazards Engineering Research Infrastructure <https://www.designsafe-ci.org/facilities/experimental/>
- [noao] National Optical Astronomy Observatory <https://www.noao.edu>
- [osg] Open Science Grid <https://opensciencegrid.org/>
- [rcrv] Regional Class Research Vessel <https://ceoas.oregonstate.edu/ships/rcrv/>
- [sgci] The Science Gateways Community Institute, <https://sciencegateways.org/>
- [trustedci] Trusted CI, the NSF Cybersecurity Center of Excellence, <https://trustedci.org/>
- [xsede] The Extreme Science and Engineering Discovery Environment (XSEDE) <https://www.xsede.org/>



Appendix A: Workshop Contributors

Steering Committee

The workshop was organized by a steering committee composed of CI experts from the LFs and selected CI projects:

- Brian Bockelman, Morgridge Institute and HTCondor project [htcondor]
- Adam Bolton, National Optical Astronomy Observatory [noao]
- Tom Cheatham, University of Utah and Campus Research Computing Consortium [carcc]
- Ewa Deelman (PI and Chair), University of Southern California and CI CoE Pilot [cicoe-pilot]
- Tom Gulbransen, Battelle, NEON [neon]
- Kate Keahey, Argonne National Laboratory, Chameleon [chameleon]
- Marina Kogan, University of Utah
- Dan Stanzione, Texas Advanced Computing Center, Leadership-Class Computing Facility [lccf]
- Daryl Swensen, Oregon State University, Regional Class Research Vessel [rcrv].

Discussion Leads

In addition to the steering committee, the following people led the breakout sessions and facilitated discussions:

- Ilya Baldin, RENCI, UNC-Chapel Hill
- Patrick Brady, University of Wisconsin-Milwaukee
- Rafael Ferreira da Silva, USC Information Sciences Institute
- Anirban Mandal, RENCI, UNC - Chapel Hill
- Jarek Nabrzyski, University of Notre Dame
- Susan Sons, Indiana University, CACR
- Alexander Szalay, Johns Hopkins
- Douglas Thain, University of Notre Dame
- John Towns, NCSA

Writing Contributors

The following people led the note taking during the meeting and contributed to this report:

- Ewa Deelman, USC Information Sciences Institute, Lead coordinating author
- Ilya Baldin, RENCI, UNC-Chapel Hill
- Laura Christopherson, RENCI, UNC-Chapel Hill
- Tom Gulbransen, Batelle
- Anirban Mandal, RENCI, UNC-Chapel Hill
- Angela Murillo, Indiana University-Purdue University Indianapolis
- Valerio Pascucci, University of Utah
- Steve Petruzza, University of Utah
- Mats Rynge, USC Information Sciences Institute
- Susan Sons, Indiana University, CACR
- Chaudhuri Surajit, Microsoft
- Charles Vardeman, University of Notre Dame
- Jane Wyngaard, University of Notre Dame



List of Attendees

First Name	Last Name	Organization
Mark	Abbot	WHOI
Rachel	Adams	UC Boulder
Stuart	Anderson	Caltech
Scot	Arnold	NSF
Ilya	Baldin	RENCI, UNC - Chapel Hill
Steve	Barnet	UW-Madison
Chaitan	Baru	NSF
Karan	Bhatia	Google
Cheryl Ann	Blain	Naval Research Laboratory
Brian	Bockelman	Morgridge Institute for Research
Adam	Bolton	NOAO
Devin	Bougie	Cornell University
Patrick	Brady	University of Wisconsin-Milwaukee
Jennifer	Bridge	University of Florida
Joel	Brock	Cornell University
Sharon	Broude Geva	University of Michigan
Peter	Bryan	Lehigh University
Robert	Chadduck	NSF
Vipin	Chaudhary	NSF
Surajit	Chaudhuri	Microsoft
Thomas	Cheatham	University of Utah
Laura	Christopherson	RENCI, UNC - Chapel Hill
Peter	Couvares	Caltech
Chris	Davis	NSF
Ewa	Deelman	USC Information Sciences Institute
Mark	Dufour	NSF/DIAS/BFA
Lee	Ellett	Scripps Institution of Oceanography, UCSD
Douglas	Ertz	UNAVCO
Ken	Feldman	University of Washington
Changda	Feng	Florida International University
Rafael	Ferreira da Silva	USC Information Sciences Institute
Douglas	Fils	Consortium for Ocean Leadership (COL)
Ian	Foster	ANL



Montona	Futrell-Griggs	NSF
Rob	Gardner	University of Chicago
Philip	Gates	Texas A & M University
Erwin	Gianchandani	NSF
Jeffrey	Glatstein	WHOI
Brian	Glendenning	NRAO
Tom	Gulbransen	Battelle Memorial Institute
David	Halstead	NRAO
Ben	Hauger	AURA, NOAO
John	Haverlack	University of Alaska
Pamela	Hill	NCAR
Margaret	Johnson	NCSA/UIUC
JJ	Kavelaars	National Research Council of Canada
Kate	Keahey	ANL
Jeff	Kern	NRAO
Marina	Kogan	University of Utah
Alexis	Lewis	NSF
Sean	Liddick	NSCL / MSU
Miron	Livny	University of Wisconsin-Madison
Eric	Lyons	University of Arizona
Shadi	Mamaghani	NSF
Anirban	Mandal	RENCI, UNC - Chapel Hill
Jasmine	Mann	USC Information Sciences Institute
Thomas	Marullo	Lehigh University
William	Miller	NSF
Angela	Murillo	Indiana University-Purdue University Indianapolis
Jarek	Nabrzyski	University of Notre Dame
Sanjay	Padhi	AWS
Manish	Parashar	NSF
Valerio	Pascucci	University of Utah
Joy	Pauschke	CMMI
Chuck	Pavloski	Penn State Institute for CyberScience
Steve	Petruzza	University of Utah
Kevin	Porter	NSF
Roland	Roberts	NSF
Stefan	Robila	NSF
Christopher	Romsos	Oregon State University, CEOAS



Jim	Rosser	IODP JRSO
Mats	Rynge	USC Information Sciences Institute
Matt	Schoettler	UC Berkeley
Susan	Sons	Indiana University, CACR
Dan	Stanzione	TACC, UT Austin
Laura	Stolp	WHOI
Alejandro	Suarez	OAC
Daryl	Swensen	Oregon State Univeristy
Alexander	Szalay	Johns Hopkins
Troy	Tanner	APL-UW
Douglas	Thain	University of Notre Dame
Christopher	Thompson	Purdue University
Kevin	Thompson	NSF
Joanne	Tornow	NSF
John	Towns	NCSA
Charles	Vardeman	University of Notre Dame
Ed	Walker	NSF
Joseph	Wartman	University of Washington
Von	Welch	CACR, Indiana University
Tim	Weston	UC Boulder
Dan	Wilson	UC Davis
Frank	Wuerthwein	UCSD/SDSC
Jane	Wyngaard	University of Notre Dame
Michael	Zentner	SDSC



Appendix B: Agenda

Day 1 – Monday, September 16, 2019

07:30 – 08:30	Breakfast and Registration
08:30 – 08:35	Opening Remarks William Miller (NSF)
08:35 – 09:05	Building LF Cyberinfrastructure Communities to Advance the Endless Frontier Joanne Tornow (NSF)
09:05 – 09:40	Setting the Stage: 2017 CI workshop and the Cyberinfrastructure Center of Excellence Pilot Ewa Deelman and Tom Gulbransen
09:40 – 10:10	Guided Activity Kate Keahay and Rafael Ferreira da Silva
10:10 – 10:40	Break
10:40 – 12:00	Panel: State and Future of Cyberinfrastructure for Large Facilities Moderator: Dan Stanzione Panelists: Stuart Anderson, Margaret Johnson, Eric Lyons
12:00 – 13:00	Lunch Break
13:00 – 13:15	NSF/CISE Perspective Erwin Gianchandani (NSF)
13:15 – 13:45	Large Facilities Data Lifecycle Anirban Mandal
13:45 – 15:15	Lightning Talks Moderator: Mats Rynge Patrick R Brady: SCIMMA: Scalable Cyberinfrastructure to support Multimessenger Astrophysics Steve Petruzza: Interactive Access and Visualization of Large Scale Image Data Miron Livny: Sustaining Software Across a Growing Number of Facilities David Halstead: Next Generation Very Large Array: Communications Jeffrey Glatstein: Analysis of Alternatives Ian Foster: Cloud CI Services for Agile Facilities Chris Romso: Software Defined Infrastructure in the Academic Research Fleet (ARF) Brian Glendenning: Scaling NRAO to the future Rob Gardner: New Approaches to Building and Operating Distributed Cyberinfrastructure for Large Facilities



	Douglas Fils: Schema.org and structured data for discovery Margaret Johnson: Enabling Multi-Instrument Pixel-Level Science with A High Throughput Computing, Data Access and Analysis Facility Kate Keahey: Chameleon: How to Build a Cloud++
15:15 - 15:30	Result Survey Overview and Setting up the Breakouts Ewa Deelman
15:30 – 16:00	Break
16:00 – 17:30	Parallel Breakouts: 1. What are the opportunities for collaboration between LFs and other Large CI projects? Leads: Ilya Baldin and Brian Bockelman 2. What are the CI challenges that need to be addressed in the next 5 years to support LFs science missions? Leads: Anirban Mandal and Daryl Swensen 3. What are the non-technical issues that influence CI development and how they can be collaboratively addressed? Leads: Susan Sons and Doug Thain
17:30 – 18:00	Breakout Summaries Top 3-5 findings and recommendations from each group
18:30 – 20:30	Reception

Day 2 – Tuesday, September 17, 2019

07:30 – 08:20	Breakfast
08:20 – 08:30	Setting the stage for Day 2 Ewa Deelman
08:30 – 10:00	Panel: Shared CI Services Opportunities and Challenges Moderator: Adam Bolton Panelists: Pamela Hill, JJ Kavelaars, Von Welch, and Mike Zentner
10:00 – 10:30	Break
10:30 – 12:00	Panel on Workforce Development and Retention Moderator: Tom Cheatham Panelists: Rachel Adams, Sharon Broude Geva, Jim Rosser, Frank Wuerthwein
12:00 – 13:00	Lunch Break



13:00 – 14:30	<p>Parallel Breakouts:</p> <ol style="list-style-type: none">1. Building a CI community: what are the impediments and opportunities? Leads: Patrick Brady and Marina Kogan2. Enhancing the CI workforce: what are the challenges and solutions? Leads: Jarek Nabrzyski and John Towns3. CI for science, where does the LF CI end and the user CI begin? Leads: Tom Gulbransen and Alex Szalay
14:30 – 15:00	<p>Breakout Summaries Top 3-5 findings and recommendations from each group</p>
15:00 – 15:15	<p>Wrap-up</p>



Appendix C: Cyberinfrastructure Calling Cards

The following CI Calling Cards have been collected from the attendees prior to the workshop.

CI accomplishment	The CONVERGE initiative added the first social science and interdisciplinary data model to the DesignSafe CI.	 <p>Rachel Adams rachel.adams-1@colorado.edu Natural Hazards Center</p>
CI frustration or challenges	Changing the data publishing culture so that social scientists use CI platforms to publish their data.	
Non-technical CI issue or success	Defining a common set of procedures and tools for storing and managing hazards and disaster data.	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

<https://converge.colorado.edu>

CI accomplishment	Worked with NEON to develop Disaster Recovery Planning process template for LFs.	 <p>Ilya Baldin ibaldin@renci.org RENCI/UNC Chapel Hill</p>
CI frustration or challenges	Custom software for LFs tends to be overlooked as a critical component of infrastructure, particularly its QA and CI/CD processes.	
Non-technical CI issue or success	Getting the right balance between planning and execution on a limited budget.	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

<https://cicoe-pilot.org>



CI accomplishment	Release of the IceProd 2 workload management system. This release simplifies using it to manage experimental data processing as well as simulation and is a step toward becoming a user facing platform.
CI frustration or challenges	Adapting a 15 year old code base to fit into contemporary CI systems and workflows.
Non-technical CI issue or success	Finding the time and resource to experiment with new computing models and analysis techniques.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Steve Barnet
barnet@icecube.wisc.edu
University of Wisconsin-Madison



IceCube Neutrino Observatory

<https://www.icecube.wisc.edu>

CI accomplishment	Using the NHERI CI web interface to connect the distributed NHERI network via an active newsroom, learning center, and Slack online collaboration.
CI frustration or challenges	Implementation of changes or additions to the CI web interface since the CI is maintained and managed by one LF in a network of distributed LFs.
Non-technical CI issue or success	The ability of the CI to connect the distributed LFs once all of the LFs began to view themselves as part of a network and not just an individual LF.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Cheryl Ann Blain
cheryl.ann.blain@nrlssc.navy.mil
Naval Research Laboratory



Natural Hazards Engineering Research Infrastructure

<https://www.designsafe-ci.org/>



CI accomplishment	Begun the process of overhauling the Authentication and Authorization Infrastructure of the OSG to adopt modern token-based technologies	 Brian Bockelman bbockelman@morgridge.org Morgridge Institute for Research
CI frustration or challenges	Shrinking the complexity of new R&D techniques to something an average PI can utilize.	
Non-technical CI issue or success	Building collaborations across projects -- without spending <i>all</i> our time building collaborations!	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure


<https://opensciencegrid.org>



CI accomplishment	Building mixed agile teams of scientists and software engineers to develop and operate online services for data-intensive astronomy research: https://datalab.noao.edu & https://antares.noao.edu	 Adam Bolton bolton@noao.edu National Optical Astronomy Observatory / AURA
CI frustration or challenges	Enabling innovation and continuous improvement within an operations context that relies on many complex legacy systems and processes	
Non-technical CI issue or success	Evolution from silo mindset to strategic mindset within organizations	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

<http://ast.noao.edu/facilities/csdc>





CI accomplishment	Continually expanding central storage cluster to handle exponential growth in data and data rates, as well as meeting new requirements for isolation and privacy.
CI frustration or challenges	No common workflow or tools for setup, data collection, and analysis across disparate beam lines and detectors. Will need to start processing and reducing data inline to keep up with ever increasing data rates.
Non-technical CI issue or success	Building momentum and starting progress on defining a common set of procedures and tools for collecting and storing data and metadata.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure




Devin Bougie
devin.bougie@cornell.edu
Cornell University

The Center for High-Energy X-ray Sciences at CHESS
www.chess.cornell.edu

CI accomplishment	Automated cyberphysical laboratory experiments which couple mechanised physical models, numerical models, and an off-site research platform in real-time.
CI frustration or challenges	Balancing experimental instrumentation development resources between functionality and security requirements within a networked environment we do not have direct control over.
Non-technical CI issue or success	Bridging the gap between NOAs sent to externally funded PIs and project onboarding.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Jennifer A. Bridge
jennifer.bridge@essie.ufl.edu
University of Florida

Natural Hazards Engineering Infrastructure @ UF
<https://ufl.designsafe-ci.org>




CI accomplishment	Complying with both Open Data Act and ITAR restrictions in same facility - meeting needs of diverse funding partners.
CI frustration or challenges	Data rates are increasing 10x per year. Currently ~10Tb/day for each user. Each experiment is unique - capturing metadata and providing visualization and/or analysis tools is critical. Experiments change almost daily on multiple beamlines.
Non-technical CI issue or success	Disconnect between CI community and diverse science community. Users need fully implemented working solutions. Generic tools are not helpful.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Joel Brock
joel.brock@cornell.edu
Cornell University



CHESS
CORNELL HIGH ENERGY SYNCHROTRON SOURCE

Cornell High Energy Synchrotron Source (CHESS)

www.chess.cornell.edu

CI accomplishment	ARC at U-M is a unique, holistic, and cohesive initiative that brings together faculty and student programmatic institutes and a training and consulting unit for CI and research computation.
CI frustration or challenges	CI is seen as infrastructure and operations, primarily, so we get bogged down in discussions about costs and ROI \$\$, rather than the really important and intangible ROI.
Non-technical CI issue or success	Professionalization, recognition, workforce development, and career paths for everyone working with, in, and supporting CI-enabled research, is hugely challenging.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Sharon Broude Geva
sgeva@umich.edu
University of Michigan



ADVANCED RESEARCH COMPUTING

arc.umich.edu



CI accomplishment	Development of a human resources wireframe and job classification for researcher-facing, systems-facing, etc. CI professionals, development of a tool to assess research computing and data maturity, and hosted a workshop bringing together the larger CI ecosystem...
CI frustration or challenges	Justifying our existence to administration and lack of sustainability in supporting people and researchers - we technical challenges, but we overcome these
Non-technical CI issue or success	Inclusivity, diversity, recognizing and embracing labels while trying to overcome them in this small CI community...

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure




Thomas Cheatham, III
tec3@utah.edu
University of Utah



CaRCC - The Campus Research Computing Consortium
carcc.org

CI accomplishment	Transitioning from a network of dedicated single-cluster HTC workflows to a Distributed HTC model, leveraging OSG to bootstrap the effort and prove its feasibility with initially limited internal resources.
CI frustration or challenges	Difficulty of funding <i>sustained</i> FTE effort for long-term collaboration between CI developers and professional project staff for ongoing, iterative CI development, integration, and support in production research environments.
Non-technical CI issue or success	Success: Improved engagement and relationships with outside CS and CI experts; looking outside first, less homegrown CI. Issue: How to support internal CI innovation (skunkworks) while limiting risk of poorly-conceived projects becoming unplanned, de-facto production services we must integrate, maintain and support.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Peter Couvares
pcouware@caltech.edu
LIGO Laboratory / Caltech



LIGO
Laser Interferometer Gravitational-Wave Observatory
<https://ligo.caltech.edu/>



CI accomplishment

The Pilot worked closely with NEON to identify working groups whose CI products can impact NEON's users

CI frustration or challenges

It is hard to find good benchmarks for evaluating scientific workflow systems

Non-technical CI issue or success

There are a lot of efforts/workshops focused on various aspects of CI, it is hard to keep track of it all

**Ewa Deelman**

deelman@isi.edu

University of Southern California

**Cyberinfrastructure Center of Excellence Pilot**<https://cicoe-pilot.org>

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

CI accomplishment


Modern, technologically advanced vessels and upcoming cybersecurity requirements are justifying improved CI. An example is the deployment of high-availability computing clusters on research vessels.

CI frustration or challenges

Lack of reliable hardened internet communications to research vessels. As a result we are limited in the effective scalable CI solutions that can be deployed at sea.


Non-technical CI issue or success

Coordination of this distributed Large Facility was established in 1972. The operational standards have evolved but have yet to define a practicable CI plan. Currently, CI competes with instrumentation and vessel operational funds, but is necessary for both.

**Lee Ellett**

lellett@ucsd.edu

Scripps Institution of Oceanography

**U. S. Academic Research Fleet**www.unols.org

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



CI accomplishment	A 24/7 data center in a smaller footprint and fewer resources.	 Douglas J Ertz dertz@unavco.org UNAVCO
CI frustration or challenges	Lots of different software systems to support/migrate/refactor with so few resources. A vendor once told me "The bible says God created the world in 7 days but he did not have an imbedded base."	
Non-technical CI issue or success	Building/keeping expertise to reduce overall development cycles. Get ahead of the curve instead of farther behind.	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

UNAVCO
Geodetic Facility for the Advancement of Geoscience (GAGE)
www.unavco.org

CI accomplishment	I am in charge of implements cyberinfrastructure and facilitates cybersecurity in Wall of Wind. I ensure testing data quality and secure data storage. I in charge of setting up telepresence during testing.	 Changda Feng cfeng@fiu.edu Florida International University
CI frustration or challenges	Manage and store large size data (~ 30 GB for each project) and high resolution video recording.	
Non-technical CI issue or success	Make sure the tested data is meaningful and store these data safely.	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

FIU NHERI Wall of Wind
Natural Hazards Engineering Infrastructure @ FIU WOW
<https://fiu.designsafe-ci.org>




CI accomplishment	Deployed high-Availability virtual computing environments at-sea for data acquisition and processing.
CI frustration or challenges	Automated internal collection and distribution of data from heterogeneous environments.
Non-technical CI issue or success	An increasing expectation and reliance on internet connected resources from an environment with very low bandwidth / not always reliable connectivity.
Non-technical CI issue or success	at-sea personnel need ongoing training/development opportunities (and time!) to effectively maintain the increasingly complex computing/networking environment on research vessels.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Ken Feldman
ksf1@uw.edu
University of Washington



U. S. Academic Research Fleet

www.unols.org

CI accomplishment	Provided open-source tools for enabling scientific progress in CI research and development
CI frustration or challenges	There is not much materials to lower the entry level barrier on how to deploy and assess CI performance metrics and reliability
Non-technical CI issue or success	Challenge to address the diversity and different levels of CI training and education

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Rafael Ferreira da Silva
rafsilva@isi.edu
University of Southern California



Cyberinfrastructure Center of Excellence Pilot

<https://cicoe-pilot.org>



CI accomplishment	Fortunate to be part of EarthCube and ESIP's interest in structured data on the web and schema.org for DataSet developments. Applying web architecture to FAIR data goals.
CI frustration or challenges	Finding the right abstraction of services and architecture that provides both ease of use and maintainability while still allowing flexible evolution and change to support emerging approaches.
Non-technical CI issue or success	Many goals pull on people & skills across many groups & organizations. However, getting organizations to free up their people's time for such tasks can be hard due to funding and reporting obligations.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Doug Fils
dfils@oceanleadership.org
Ocean Leadership



Project 418/419

gleaner.io

CI accomplishment	Globus data, identity, and automation services are used by >1000 institutions and 100,000 people worldwide. Large facilities such as NCAR, XSEDE, and DKIST simplify operations by outsourcing to Globus services.
CI frustration or challenges	Accelerating progress towards a global science platform to enable collaboration and sharing at scale.
Non-technical CI issue or success	Realization of a successful sustainability solution, based on subscriptions from institutions and facilities, while preserving free access to core capabilities.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Ian Foster
foster@uchicago.edu
UChicago, Argonne, Globus



**Computer Science, UChicago;
Data Science & Learning,
Argonne;
Globus, UChicago**
www.globus.org



CI accomplishment	Supporting global cyberinfrastructure for frontier science including LHC, dark matter searches and cosmic microwave background studies	 <p>Rob Gardner rwg@uchicago.edu University of Chicago</p>
CI frustration or challenges	Integration of new and heterogeneous resources into shared, collaborative computing environments	
Non-technical CI issue or success	Development of a new trust model for operation of distributed service platforms for multi-institution collaborations	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Services Layer at the Edge

slateci.io

CI accomplishment	Our team worked with Texas A&M University Infrastructure Operations to execute a two-day migration of significant IODP servers and storage to their new, state-of-the-art data center.	 <p>Philip Gates gates@iodp.tamu.edu Texas A&M University</p>
CI frustration or challenges	Providing safe, reliable shipboard network and application access for a multinational science party with a diverse assortment of laptops, operating systems and software versions.	
Non-technical CI issue or success	A Project Management Office with well-defined guidelines for project management, planning, and reporting brings communication, predictability, transparency, and efficiency.	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



**JOIDES Resolution
Science Operator
(IODP)**

<https://iodp.tamu.edu/>



CI accomplishment	Implementation of monitoring and metrics which allow reporting to stakeholders on system use, leading to both system stability and more transparency with regard to system status.
CI frustration or challenges	Data Portal user experience challenges serving up complex technical data
Non-technical CI issue or success	Gaining understanding and consensus that CI implementations need to continually grow and reassess vs have a set completion date with full step down to minimal staffing and budget.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Jeffrey Glatstein
jglatstein@whoi.edu
Woods Hole Oceanographic Institution



Ocean Observatories Initiative

<https://oceanobservatories.org>

CI accomplishment	HTCondor / OSG being adopted as good architectural fit for out-of-institute computing resources needed for various Observatory initiatives
CI frustration or challenges	Parts of codebase for main data processing package ~25 years old; parallelization/performance and maintenance challenges
Non-technical CI issue or success	Hiring software development staff is currently difficult (industry salary scales, location, old technology)

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Brian Glendenning
bglenden@nrao.edu
National Radio Astronomy Observatory



www.nrao.edu



CI accomplishment	Optimizing completion of NEON cyberinfrastructure construction and launching operations within schedule and budget constraints.
CI frustration or challenges	Need for more Data Scientists who blend domain knowledge in ecological sciences, with quantitative analytical methods, and computer science experience.
Non-technical CI issue or success	Increasing power of using semantics to strengthen data discovery & integration in workflow-driven analyses.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Tom Gulbransen
gulbransen@battelle.org
Battelle - NEON



The National Ecological Observatory Network (NEON)
<https://www.neonscience.org>

CI accomplishment	Provisioning 10Gbps capable fiber connectivity to 6 radio telescopes in remote locations
CI frustration or challenges	No cost-effective way of engaging a single entity for coordinating the delivery of “last mile” fiber infrastructure for access to Internet 2 backbone
Non-technical CI issue or success	The cadence of large Radio Telescope projects is typically 15-20 years from proposal to first light. Partnership with National Computing Centers is a challenge due to their ~5 year CI funding horizons

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure




David M. Halstead
dhalstead(at)nrao.edu
NRAO



National Radio Astronomy Observatory
www.nrao.edu



CI accomplishment	<p>Documentation with MkDocs + Git</p> <ul style="list-style-type: none">• https://wiki.cfos.uaf.edu/docs/ <p>Documentation renders as searchable static HTML + PDF, is cross platform, uses markdown files for content, can easily be hosted securely on websites (No back end database). Content can be fully version managed in Git, and most importantly you can have full searchable local access to documentation on a thumb drive or laptop without access to online resources.</p>	 <p>John Haverlack jehaverlack@alaska.edu University of Alaska Fairbanks</p>
CI frustration or challenges	<p>We constantly struggle for depth of team resources for:</p> <ul style="list-style-type: none">• Responsibility Matrix Cross Service Coverage• Development to Production Pipeline• Keeping up with fast paced changing technology• Technical Debt	
Non-technical CI issue or success	<ul style="list-style-type: none">• Consistent Service Management Practices (e.g. ITIL)• Business Continuity• Shared Collaboration Tools, Processes and Practices• Herding Cats	
2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure		https://www.unols.org/



U.S. Academic Research Fleet

CI accomplishment	<p>Implementation of a new 'Campaign Store' resource for med-term data archival (1-5 years).</p> <ul style="list-style-type: none">• Initial deployment of 26 PB• Annual increments between 15-20 PB	 <p>Pamela Hill pjg@ucar.edu NCAR</p>
CI frustration or challenges	<p>Lack of data management practices and tools</p> <ul style="list-style-type: none">• we currently have ~2 billion files across active, campaign and archival storage• no one knows what they have or how to begin to unravel	
Non-technical CI issue or success	<p>Lack of understanding how new technologies do/do not fit our users workflows</p> <ul style="list-style-type: none">• data formats don't work with object storage solutions• data transfer rates to cloud are too slow	
2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure		Project URL



HPC Storage Architectures



CI accomplishment	Created an architectural framework for developing a strategic plan for complex IT facilities, based on enterprise architecture concepts. Applied for the LSST Data Facility, proposed for use in other contexts.
CI frustration or challenges	How to ensure operational thinking gets integrated early on in projects, and keeping that frame through design and construction.
Non-technical CI issue or success	<ul style="list-style-type: none">● Resistance to leverage industry-standard concepts and frameworks (e.g., enterprise architecture, IT service management, IT financial management)● Tendency to choose of tools in lieu of building processes

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Margaret Johnson
mgelman2@illinois.edu
NCSA



LSST
Large Synoptic Survey Telescope

<https://www.lsst.org>

CI accomplishment	Integration of all metadata search across all data collections and inclusion of metadata from external CI sites meta-data within those search systems.
CI frustration or challenges	Deployment of new systems involves interactions across 3 different 'shared service' organizations and getting their attention to our problems hard.
Non-technical CI issue or success	Brought 'boots on the ground' to team-launch meeting to establish personal connection: communication and commitment increased dramatically

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure




JJ Kavelaars
jj.kavelaars@nrc.gc.ca
National Research Council of Canada




Canadian Astronomy Data Centre
www.cadc.hia.nrc.gc.ca



CI accomplishment	<p>Created and operates Chameleon, a testbed for Computer Science research based on mainstream cloud technology</p>	 <p>Kate Keahey keahay@anl.gov Argonne Nat'l Lab / University of Chicago</p>  <p>Chameleon</p>
CI frustration or challenges	<p>Creating a sharing/publishing workflow for replicable experiments</p>	
Non-technical CI issue or success	<p>Creating and managing incentives for fair sharing in non-monetary economies is hard to yet increasingly critical to ensure as the needs for sharing different resources in different ways evolve</p>	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

www.chameleoncloud.org

CI accomplishment	<p>In the process of creating a rich archive for radio astronomy that supports data discovery, analysis, and reprocessing for the astronomy community.</p>	 <p>Jeff Kern jkern@nrao.edu NRAO</p>  <p>National Radio Astronomy Observatory</p>
CI frustration or challenges	<p>Compute requirements for VLA Sky Survey are at least an order of magnitude larger than predicted. We need to engage with CI partners to deliver our science.</p>	
Non-technical CI issue or success	<p>Finding the right “fit” of CI resources and approach to our problem has been a challenge.</p>	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

www.nrao.edu



CI accomplishment	Led development of an increase in front-end data collection throughput by an order of magnitude
CI frustration or challenges	Sharing large data sets from the facility with geographically distributed end users
Non-technical CI issue or success	Lack of a coordinated framework for consolidating, evaluating, and disseminating best practices for CI. Implemented ISO27001 within facility.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Sean Liddick
liddick@nscl.msu.edu
Michigan State University



National Superconducting Cyclotron Laboratory
www.nscl.msu.edu

CI accomplishment	NRAO interest in transitioning their computing to an HTC model and adopting HTCondor and OSG to support their Capabilities, Workflows and Tasks
CI frustration or challenges	Manage the allocation of (distributed) storage
Non-technical CI issue or success	The little interest of the community in understanding principles and the fascination with buzz words and hype

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Miron Livny
miron@cs.wisc.edu
University of Wisconsin-Madison



UW-Madison Center For High Throughput Computing
<http://chtc.cs.wisc.edu/>



CI accomplishment	Over 66,000 users worldwide; 3.2PBs of shared research data (99% is private) used by researchers (7PB annual traffic); launch of VICE for real-time interactive computing using Jupyter Notebooks and R Shiny Apps
CI frustration or challenges	CI doesn't automatically create documentation that is useful to all users (regardless of their expertise and background); no matter how good people become and creating software, it still sucks to get it integrated in the CI
Non-technical CI issue or success	Scaling people -- major challenge in addressing the large diversity of training needs for all types of researchers

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Eric Lyons
ericlyons@email.arizona.edu
University of Arizona



CyVerse
<https://cyverse.org>

CI accomplishment	Worked closely with NEON and IceCube Large Facilities to develop a Data Life Cycle abstraction and Disaster Recovery Planning Effective Processes for Large Facilities.
CI frustration or challenges	Complexity of Data Life Cycle and heterogeneity of underlying services, architectures and tools for Large Facility CI.
Non-technical CI issue or success	Cataloging best practices and catalyzing communication across common CI issues in Large Facilities are important, but missing in the current ecosystem.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Anirban Mandal
anirban@renci.org
RENCI, UNC - Chapel Hill



**Cyberinfrastructure
Center of Excellence Pilot**
<https://cicoe-pilot.org>



CI accomplishment	Developed IT infrastructure at Lehigh University ATLSS Center for the NHERI Lehigh program which synchronizes with NHERI DesignSafe-CI
CI frustration or challenges	Interacting with off-site data repositories or physical test systems in real-time
Non-technical CI issue or success	Freedom to implement new ideas or technologies

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Thomas Marullo


tmm3@lehigh.edu

Lehigh University



CI accomplishment	Examining various aspects of DataONE CI to assist with ensuring data sharing and reusability.
CI frustration or challenges	Concerns with data management and curation practices.
Non-technical CI issue or success	Rapidly changing technologies, the number of projects and potential solutions.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Angela Murillo

apmurill@iu.edu

Indiana University-Purdue University Indianapolis



Indiana University-Indianapolis

soic.iupui.edu



CI accomplishment	Led building a complex CI development and support team (a center) at a mid-size private University through hiring and developing talent from 18 countries worldwide (total ~60 staff and faculty)
CI frustration or challenges	Working on proposals, grants and contracts to support over 70% of all salaries in the center
Non-technical CI issue or success	Competing internally with other University priorities that change on an annual basis. Lack of effective research IT governance models across colleges, central IT and administration.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Jarek Nabrzyski
naber@nd.edu
University of Notre Dame



Cyberinfrastructure Center of Excellence Pilot

<https://cicoe-pilot.org>

CI accomplishment	Enabled interactive visual exploration of Airborne Data (multispectral photography) of the NSF National Ecological Observatory Network (NEON) data portal
CI frustration or challenges	Help the community develop integrated Data Management, Analysis and Visualization solutions for Big Data that are leveraged across different applications
Non-technical CI issue or success	Balance the quick deployment of new technical solutions to address existing challenges with the need for the users to work with stable and reliable solutions

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Valerio Pascucci
pascucci@sci.utah.edu
University of Utah




Cyberinfrastructure Center of Excellence Pilot

<https://cicoe-pilot.org>



CI accomplishment	We recently expanded our system offerings to include ICS-ACI's HPRC (High Performance Research Cloud) which pushed our facility core count to 36k and enabled our virtual core / virtual machines for HTC.
CI frustration or challenges	Our biggest challenge is making sure we are prepared for the future needs of our research faculty with respect to our data center and power needs
Non-technical CI issue or success	<ol style="list-style-type: none">1 - Making sure the engineering and security process does not slow deployment of new offerings2 - Convincing leadership to allow us to focus on research computing and not other IT related issues

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Chuck Pavloski
cfp102@psu.edu
Penn State



Penn State Institute for
CyberScience Advanced
CyberInfrastructure System

<https://ics.psu.edu>

CI accomplishment	Deployed infrastructure for interactive exploration of Airborne Observatory Platform (AOP) Data and coordinate integration with the NSF National Ecological Observatory Network (NEON) data portal.
CI frustration or challenges	Lack of interactive data access limits the ability of users to discover and retrieve the data they need in a timely manner.
Non-technical CI issue or success	Open source software and stable and usable deployments are essential for fast adoption and extension of existing software solutions.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Steve Petruzza
spetruzza@sci.utah.edu
University of Utah




Cyberinfrastructure
Center of Excellence Pilot

<https://cicoe-pilot.org>



CI accomplishment	<p>Technical lead for the development of the Cruise Observations Real-time Interface & Open Live Information eXchange (CORIOLIX). An application for bidirectional (ship <-> shore) data exchange.</p>	 <p>Chris Romsos chris.romsos@oregonstate.edu Oregon State University, CEOAS</p>  <p>Regional Class Research Vessel Project</p> <p>https://ceoas.oregonstate.edu/ships/rcrv/</p>
CI frustration or challenges	<p>Biggest challenge? That's easy. Delivering robust, reliable, and standardized CI for the 3 new vessels into a heterogeneous academic research fleet with independent vessel operators.</p>	
Non-technical CI issue or success	<p>It's probably a bit early to declare success. A non-technical problem on the horizon will be training our marine technical teams to maintain and operate our shipboard and shoreside CI.</p>	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

CI accomplishment	<p>Our CI team successfully migrated all shore-based data services to Texas A&M University's new West Campus Data Center in two days as part of a university-wide shared services initiative.</p>	 <p>Jim Rosser jrosser@tamu.edu Texas A&M University</p>  <p>JOIDES Resolution Science Operator (IODP)</p> <p>https://iodp.tamu.edu</p>
CI frustration or challenges	<p>Complying with NIST-based cybersecurity controls without negatively impacting customers and mission accomplishment.</p>	
Non-technical CI issue or success	<p>We employ an outstanding team of technologists who are passionate about providing our customers with exceptional service.</p>	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



CI accomplishment	The Pilot worked closely with NEON to identify working groups whose CI products can impact NEON's users
CI frustration or challenges	CI is a never-ending moving target, with many competing standards
Non-technical CI issue or success	Developing and retaining a great CI workforce.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Mats Rynge
rynge@isi.edu
University of Southern California



Cyberinfrastructure Center of Excellence Pilot

<https://cicoe-pilot.org>

CI accomplishment	Deployment of regional hazard simulation workflows at DesignSafe-CI (TACC) that are demonstrated through large-scale testbeds for seismic (Anchorage, AK and San Francisco Bay Area) and hurricane (Atlantic City, NJ) events.
CI frustration or challenges	Accessing meaningful and useful data or models published by the community.
Non-technical CI issue or success	A competitive job market makes recruitment of technical staff a challenge in a strong tech environment.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Matt Schoettler
schoettler@berkeley.edu
UC Berkeley



SimCenter NHERI
Center for Computational Modeling and Simulation

Computational Modeling and Simulation Center

simcenter.designsafe-ci.org



CI accomplishment	<p>Helping NEON to adopt a new IdM solution and mature their IdM practice more generally.</p> <p>Standing up the ResearchSOC service.</p>	 <p>Susan Sons sesons@iu.edu Indiana University</p>
CI frustration or challenges	<p>A culture of building infrastructure but not maintaining it...which results in either constant migration or constant brokenness, or sometimes balancing both types of pain.</p>	
Non-technical CI issue or success	<p>Attracting, developing, and keeping a highly-skilled workforce in a world where the private sector pays more, hires faster, and has less credentialism</p>	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure


CI CoE Pilot & ResearchSOC
<https://cicoe-pilot.org>
<https://researchsoc.iu.edu>

CI accomplishment	<p>Frontera - The largest academic supercomputer in the world - and the project to support science that goes with it.</p>	 <p>Dan Stanzione dan@tacc.utexas.edu Texas Advanced Computing Center/LIT-Austin</p>
CI frustration or challenges	<p>The fragmentation of the marketplace - software, tools, languages, workforce skills, etc.</p>	
Non-technical CI issue or success	<p>We've helped several other people win Nobel prizes -- and thousands of others make real advances. What can be more rewarding than that?</p>	

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure


TACC
TEXAS ADVANCED COMPUTING CENTER
www.tacc.utexas.edu



CI accomplishment

Every cruise dataset collected is a CI accomplishment at so many levels.

CI frustration or challenges

Keeping computer systems patched and secure while they are traveling the world oceans collecting data.

Non-technical CI issue or success

There are so many aspects to data collection; actual collection, calibration, best practices, QA, DP, distribution with proper metadata, security, backups, workflows, documentation, collaboration

**Laura Stolp**

lstolp@whoi.edu

WHOI

**WHOI SSSG/ R2R**<http://rvdata.us>

CI accomplishment

The RCRV is still early in our CI development and procurement. Instituting efficient and thorough procurement process to ensure we are properly justifying and receiving the CI we need.

CI frustration or challenges

The proper gauging of what we want for modern CI and the continuation of proven technology. Integrating modern ship operations CI needs with existing known science CI system.

Non-technical CI issue or success

Working with multiple and sometimes conflicting interests for the design of our vessel CI.

**Daryl Swensen**

daryl.swensen@oregonstate.edu

Oregon State University, CEOAS

**Regional Class Research Vessel Project**<https://ceoas.oregonstate.edu/ships/rcrv/>



CI accomplishment

We have built the SkyServer for the Sloan Digital Sky Survey data, and then modified it as part of the DIBBS program to accommodate data from many different science domains. This became the SciServer.

CI frustration or challenges

How we cannot deal with the problem of the “long tail” data.

Non-technical CI issue or success

We saw 7M non-astronomers using the SDSS data from the SkyServer and GalaxyZoo.

**Alex Szalay**

szalay@jhu.edu

The Johns Hopkins University

SciServer

CI accomplishment

Development of RApp, an iPad based app for disaster reconnaissance. Data collected in RApp is automatically tagged and uploaded to DesignSafe, a cloud-based data repository.

CI frustration or challenges

It is a challenge to get managers to understand that security, testing, maintenance, and updates add cost and time to development efforts.

Non-technical CI issue or success

Success: Positive feedback from user communities

**Troy Tanner**

troyt@apl.uw.edu

University of Washington

**NHERI RAPID Facility**



CI accomplishment


Enabled portability of complex species-modeling workflows across multiple facilities via dynamic partitioning using Makeflow + JX.

CI frustration or challenges

Dealing with dependency management in software and facilities; how can we help users to better understand the stack of what they rely upon?

Non-technical CI issue or success

Finding the right time and place to reach out to potential user communities: how to make contact when searching for solutions but not yet committed.

**Douglas Thain**

dthain@nd.edu

University of Notre Dame



2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure

<http://ccl.cse.nd.edu>

CI accomplishment

As the coordination office for the NHERI program, our office is not involved directly with challenges of operating CI equipment. We hope to provide whatever support we can to CI issues throughout the program network with governance and intra-facility coordination efforts.

CI frustration or challenges

As the coordination office for the NHERI program, our office is not involved directly with challenges of operating CI equipment. We hope to provide whatever support we can to CI issues throughout the program network with governance and intra-facility coordination efforts.

Non-technical CI issue or success

In addition to maintaining the governance groups of the program, our office also oversees the cross-domain community-wide communication efforts and has to find balance highlighting each separate member awardee in news and releases without overwhelming the users or losing sight of the major accomplishments of any particular facility.

**Christopher Thompson**

thompscs@purdue.edu

Purdue University



NCO **NHERI**
Network Coordination Office
Natural Hazards Engineering
Research Infrastructure --
Network Coordination Office
(NCO)

<https://www.designsafe-ci.org>

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



CI accomplishment

Helping bring together NEON and NCAR to discuss how to enable the community to leverage their data sets and resources to the fullest.

CI frustration or challenges

Finding feasible ways to connect the rich but diverse data sets available to the biological sciences research community.

Non-technical CI issue or success

A big challenge in building CI is building a diverse and inclusive CI-savvy workforce.

**Joanne Tornow**jtornow@nsf.gov

National Science Foundation

nsf.gov

CI accomplishment

Creation of a highly effective and efficient CI supporting a broad range of researchers via infrastructure, services, support, workforce development and technical expertise.

CI frustration or challenges

Lack of a national roadmap/blueprint/plan to guide the investments made by federal agencies and others to maximize their impact on the advancement of science.

Non-technical CI issue or success

Success: Synergizing the efforts of a large number of participants in realizing the XSEDE distributed environment.
Frustration: Parochialism that prevents the leveraging of common/shared CI infrastructure and services.

**John Towns**jtowns@ncsa.illinois.edu

NCSA/University of Illinois

**XSEDE**<http://xsede.org>



CI accomplishment	The Pilot worked closely with NEON through ESIP and in collaboration with other collaborative projects to implement schema.org dataset markup
CI frustration or challenges	Connecting Semantic and Linked Data technologies to research programmers and web developers as a practical tool
Non-technical CI issue or success	Expanding community connections through ESIP and RDA between NEON and other projects with the same set of needs for semantics and schema.org markup

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Charles F Vardeman II
cvardema@nd.edu
University of Notre Dame



Cyberinfrastructure Center of Excellence Pilot

<https://cicoe-pilot.org>

CI accomplishment	Ongoing engagement with all 23 NSF Major Facilities through Large Facilities Security Team. Deep ongoing engagements with US ARF for cybersecurity program design and NRAO on cybersecurity operations.
CI frustration or challenges	The unending stream of technology silver bullets.
Non-technical CI issue or success	Project management skills.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Von Welch
vwelch@iu.edu
Indiana University



Trusted CI
THE NSF CYBERSECURITY CENTER OF EXCELLENCE
ResearchSOC

Trusted CI and ResearchSOC

trustedci.org / researchsoc.iu.edu



CI accomplishment	Recently implemented a new firewall, virtual desktop environment, and user authentication protocols, to protect networked controlled equipment used by local and remote users. Changes were implemented as we transitioned IT support from facility-dedicated to college shared services personnel.
CI frustration or challenges	As part of a national shared network of physical resources, researchers, equipment operators, and cyber infrastructure operate under discrete awards with independent deliverables. It remains frustratingly obscure to find the data produced at shared facilities.
Non-technical CI issue or success	Our biggest CI success under NEES (predecessor to NHERI) was to change the research culture in earthquake engineering from “why would I share my experimental data” to “why would I not share my experimental data.” My biggest frustration is how challenging it remains to share data even when people want to do it.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Dan Wilson
dxwilson@ucdavis.edu
UC Davis

UCDAVIS
CENTER FOR GEOTECHNICAL MODELING
NHERI 
Natural Hazards Engineering Research Infrastructure

NHERI Centrifuge facility

cgm.engr.ucdavis.edu

CI accomplishment	Developing, deploying, and operating Data Federations for CMS@LHC and Open Science at large. This includes caches in the network backbone and at clusters across the globe.
CI frustration or challenges	The biggest challenge in CMS@LHC is to transition the experiment to software and computing broadly defined that does not blow up the budget. We are talking about Exabyte/year of new data starting ~2027.
Non-technical CI issue or success	Building, deploying, and operating CI for all of open science at all scales. From individuals to 1000++ global collaborations. Including institutions at all scales and all domains, museums, colleges, national labs. What can they share and what can't they share, and why? There are often more social/political than technical challenges.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Frank Wuerthwein
fkw@ucsd.edu
UCSD/SDSC

CMS and OSG

opensciencegrid.org
[CMS](http://cms.cern.ch)



CI accomplishment	Regarding processes and technological stacks used in sensor front end data capture and initial transport, the Pilot worked with NEON and other LFs to determine common requirements and missing components.
CI frustration or challenges	Lack of best practices and mature open source tooling for large scale deployment of edge compute supported sensing systems given the rapid rate of change in this technological sphere.
Non-technical CI issue or success	The non-agile nature of LF and consequent limited ability to respond to evolving technologies due to the pragmatic constraints of research funding timelines, and budget period constrained developers.

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Jane Wyngaard
jwyngaard@nd.edu
University of Notre Dame



Cyberinfrastructure Center of Excellence Pilot

<https://cicoe-pilot.org>

CI accomplishment	Assumed leadership of the Science Gateways Community Institute, moved HUBzero to SDSC, operating nanoHUB from SDSC serving ~2M users across all activities
CI frustration or challenges	Technical debt - inertia caused by having a program that has been alive for nearly two decades
Non-technical CI issue or success	Funded efforts that are re-inventing the wheel, a lack of an overall picture of what fits where and why we should not be re-developing it again

2019 NSF Workshop on Connecting Large Facilities and Cyberinfrastructure



Michael Zentner
mzentner@ucsd.edu
San Diego Supercomputer Center



SGCI
Science Gateways Community Institute
Science Gateways Community Institute, HUBzero, nanoHUB
sciencegateways.org